

Lecture 6. CPU Microarchitecture

Introductions to Data Systems and Data Design

Ce Zhang

CPU Microarchitecture Overview

What is a CPU?

At its core, a CPU does two things:

1. **Process instructions** (compute)
2. **Move data** (memory access)

What is a CPU?

At its core, a CPU does two things:

1. **Process instructions** (compute)
2. **Move data** (memory access)

Key insight: Modern CPUs are very good at (1), but fundamentally limited by (2).

Understanding *why* will help you design better data systems.

Skylake Server Microarchitecture

Front End

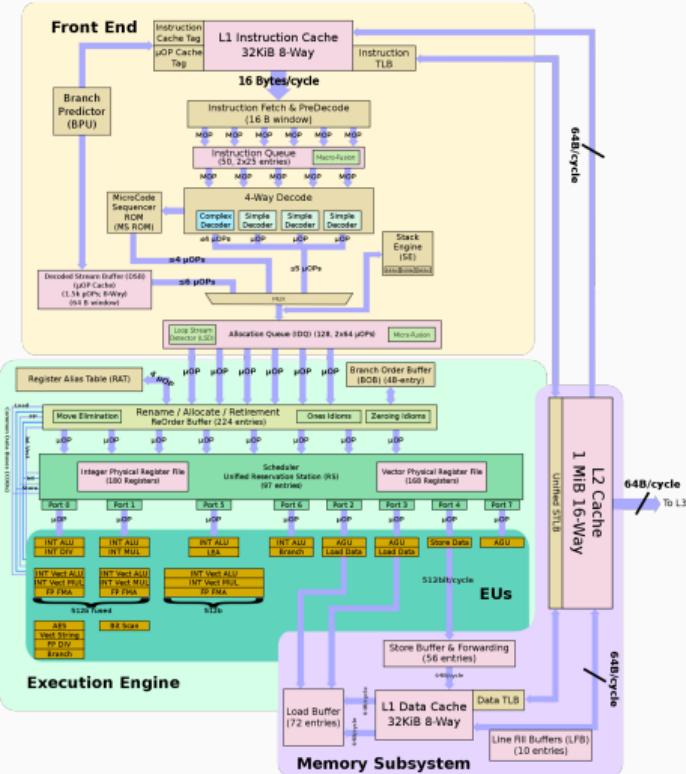
- Fetch instructions at 16 B/cycle
- Decode them into uOps

Execution Engine

- 8 ports, 4 ALUs

Memory Subsystem

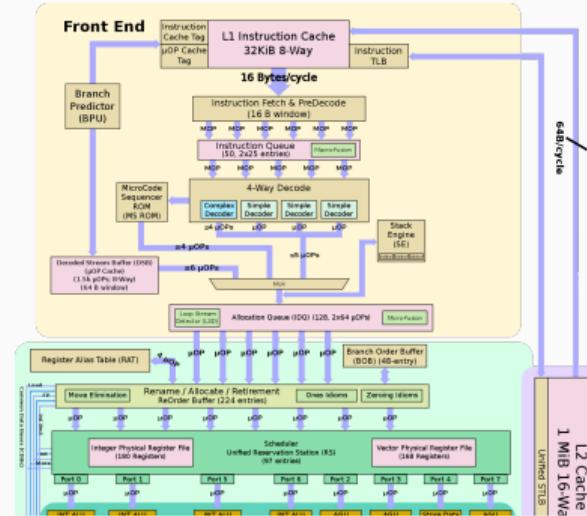
- L1D: 32 KB, 8-way, **4 cycles**
- L2: 1 MB, 16-way, **14 cycles**
- Load buffer: 72 entries
- Store buffer: 56 entries
- 64 B cache line, 64 B/cycle bandwidth



The Front End (Brief Overview)

What it does: Fetches and decodes x86 instructions into μOPs

- **L1I cache (32 KB, 8-way):** Stores instruction bytes. 8-way = each address can map to 8 cache slots (reduces conflicts).
- **Fetch (16 B/cycle):** Reads from L1I using program counter (PC) → outputs raw instruction bytes.
- **Decode (4 instr/cycle):** Translates x86 bytes → μOPs. (1 complex + 3 simple decoders)
- **μOP cache (6 μOPs/cycle):** Caches decoded μOPs for hot loops — skips decode entirely.



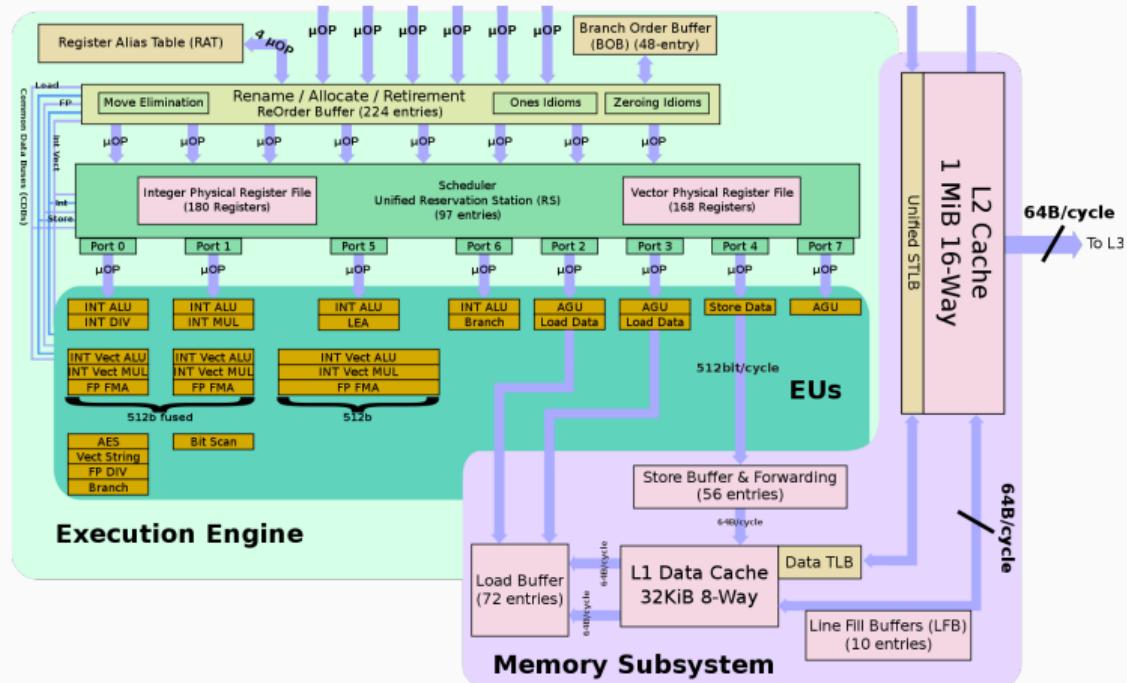
Front End Bottlenecks

When does the front end stall?

- **I-cache miss:** Code doesn't fit in 32 KB L1I → wait ~14 cycles (L2) or ~100+ cycles (RAM)
 - *Example:* Large binaries, heavy inlining, many cold function calls
- **Branch mispredict:** CPU guesses branch direction *before* condition is evaluated (to keep pipeline full). Wrong guess → flush all speculative work, restart.
 - *Predictable:* Loop branches (`for i < n`), always-taken/never-taken → >99% accuracy
 - *Unpredictable:* Data-dependent branches like `if (array[i] > threshold)` on random data → ~50% accuracy
 - ~15–20 cycle penalty per mispredict

I-cache is real! We will not be focusing on the front-end in this course, but they will show up (often unexpectedly).

Our Focus: The Execution Engine & Memory



Execution Engine (yellow) + Memory Subsystem (green) + L2 Cache (right)

The Scheduler

What it does: Holds decoded μOPs waiting for operands, dispatches to execution ports when ready.

- **97 entries** — can hold ~97 in-flight μOPs
- **Out-of-order execution:** Instructions don't execute in program order; they execute when *operands are ready*
- **Dependency tracking:** Tracks which μOPs depend on which results

Key insight: The scheduler hides latency by doing *other work* while waiting for slow operations (e.g., memory loads).

Registers and Register Renaming

Architectural registers: What the programmer sees

- **16 general-purpose** (64-bit each): rax, rbx, rcx, rdx, rsi, rdi, rbp, rsp, r8–r15
 - Can access as 32-bit (eax), 16-bit (ax), or 8-bit (al, ah)
 - Total: 16×8 bytes = 128 bytes
- **16/32 vector registers** (AVX-256: 256-bit each, AVX-512: 512-bit): ymm0–ymm15 or zmm0–zmm31
 - Used for SIMD: process 4 doubles or 8 floats in one instruction
 - Total: 16×32 bytes = 512 bytes (AVX-256)

Physical registers: What the hardware actually has (Skylake: 180 integer, 168 vector). We will not go into details, do ask ChatGPT if you are curious!

Register access: ~0 cycles latency (same cycle), “unlimited bandwidth”

Execution Ports (Skylake)

How it works:

1. Scheduler checks which μOPs have all operands ready
2. Each cycle, scheduler dispatches ready μOPs to available ports
3. Each port has specific execution units — μOP goes to a port that can handle it
4. Execution takes 1+ cycles depending on operation (add: 1 cycle, divide: 10+ cycles)
5. Result is written back, waking up dependent μOPs

Key constraint: Each port can only execute **one μOP per cycle**

Port	What it does	Latency
P0, P1	ALU, FMA (integer/FP math)	1 / 4 cycles
P0 only	Division	10-20 cycles
P5	ALU, vector shuffle	1 cycle
P6	ALU, branches	1 cycle
P2, P3	Load (2 loads/cycle)	4-5 cycles (L1)
P4, P7	Store data + address	~4 cycles

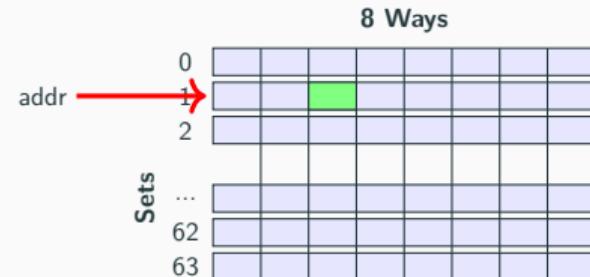
L1 Data Cache (32 KB, 8-Way Set Associative)

$$32 \text{ KB} = 64 \text{ sets} \times 8 \text{ ways} \times 64 \text{ B}$$

- **Cache line:** 64 bytes — minimum transfer (even for 1 byte read)
- **Latency:** 4-5 cycles

Mental model for this course:

- 32 KB of fast memory
- 64 B cache line (minimum unit)
- ~**LRU eviction:** Least recently used data gets evicted when full
 - It is not real LRU for two reasons: (1) 8-way set associative and (2) not exact LRU implementation.
 - In this course we will ignore these details and only think about L1 as one big chunk of cache.



L2 Cache (1 MB, 16-Way Set Associative)

1 MB per core, private

- **Cache line:** 64 bytes
- **Latency:** 14 cycles (3.5× slower than L1)

Mental model for this course:

- 1 MB of medium-speed memory per core
- **~LRU eviction**
- **Unified:** Holds both instructions and data
- Miss L1 → check L2 → 14 cycle penalty

L3 Cache (Last Level Cache)

Size: 1.375 MB per core (shared across all cores)

Latency: 50–70 cycles

Bandwidth: Lower than L2, shared among cores

Key properties:

- **Shared:** All cores access the same L3
- **Coherent:** Hardware keeps data consistent across cores

When you miss L3: DRAM access — 100+ cycles, ~100 ns.

Putting All Numbers Together

Operation	Latency	Relative
Register access	~0.3 ns	1x
L1 cache hit	~1 ns	3x
L2 cache hit	~4 ns	13x
L3 cache hit	~12 ns	40x
Main memory	~100 ns	333x
SSD random read	~16,000 ns	53,000x
HDD random read	~2,000,000 ns	6.6M x

Memory access is 100–300× slower than computation! This is *the* fundamental theme that has governed data systems design and optimization for decades.

Life Cycle of an Instruction

What Happens When You Run Code?

Every instruction goes through a journey:



What Happens When You Run Code?

Every instruction goes through a journey:



A Simple C++ Program

```
// add.cpp          $ cd examples && make add      # Compile without optimization
int main() {       $ ./add && echo $?           # Execute, print return value
    long a = 5;
    long b = 3;        $ make asm                 # Generate assembly files
    long c = a + b;   $ cat add.s                  # View assembly
    return c;         $ objdump -d add | less     # Or disassemble binary
}
```

Try it! Use godbolt.org to see assembly instantly in browser.

The Assembly (Unoptimized)

main:

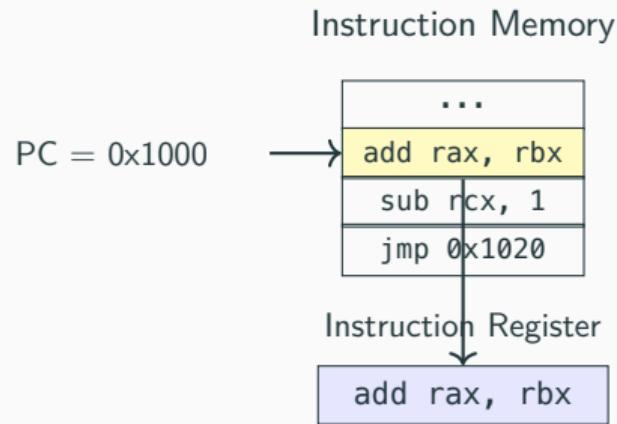
```
push rbp ; Save caller's base pointer
mov rbp, rsp ; Set up our stack frame

mov QWORD PTR [rbp-8], 5 ; Store 5 on stack (variable a)
mov QWORD PTR [rbp-16], 3 ; Store 3 on stack (variable b)

mov rdx, QWORD PTR [rbp-8] ; Load a (5) into rdx
mov rax, QWORD PTR [rbp-16] ; Load b (3) into rax
add rax, rdx ; rax = rax + rdx = 3 + 5 = 8

mov QWORD PTR [rbp-24], rax ; Store result on stack (variable c)
mov rax, QWORD PTR [rbp-24] ; Load c into rax (return value)
```

Stage 1: Fetch



- Read instruction bytes from memory (via I-cache)

Stage 2: Decode

"add rax, rbx" (bytes: 48 01 D8)



Decoder

Opcode: ADD

Source 1: RAX

Source 2: RBX

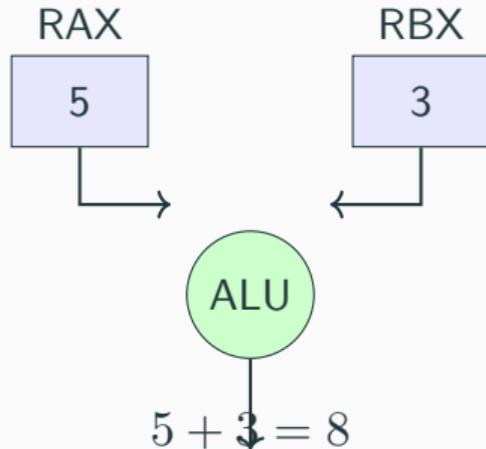
Dest: RAX



μops to execution

Modern CPUs decode complex instructions into simpler micro-operations.

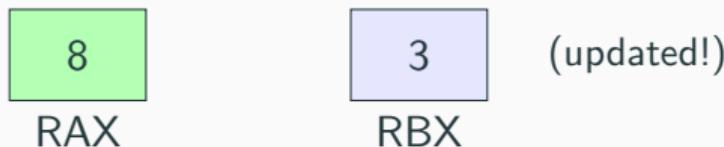
Stage 3: Execute



- Read operands from registers
- Perform the computation
- ALU operations are **fast**: 1 cycle for this case

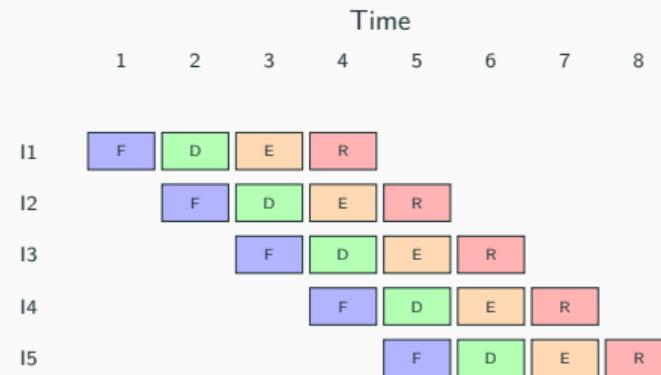
Stage 4: Retire (Write Back)

Result: 8



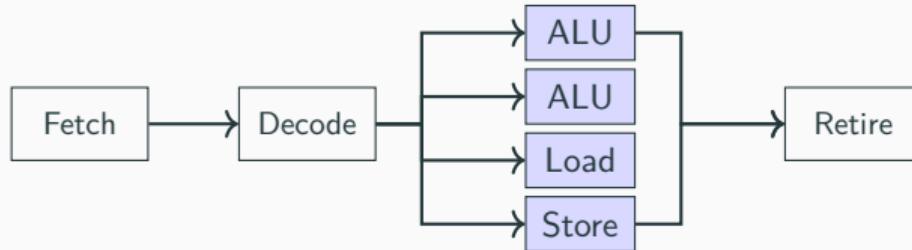
- Write result back to destination register
- Update architectural state
- Instruction is now “complete”

Pipelining: Overlapping Instructions



- After pipeline fills: 1 instruction completes per cycle
- Modern CPUs: 14-20 stages, 4-6 instructions/cycle

Superscalar: Instruction-Level Parallelism



Multiple instructions execute **in parallel** each cycle.

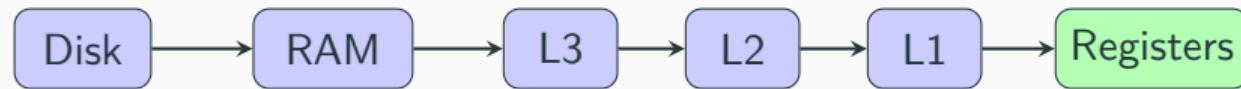
All modern CPUs are superscalar. Intel, AMD, Apple Silicon, ARM — they all execute multiple instructions per cycle.

Mental model for this course: Think of the CPU as being able to execute $\sim K$ independent instructions per cycle. If instruction B depends on the result of instruction A, they must run sequentially. If they are independent, they can run in parallel, as long as there are available *ports that can run the operation*.

Life Cycle of Data

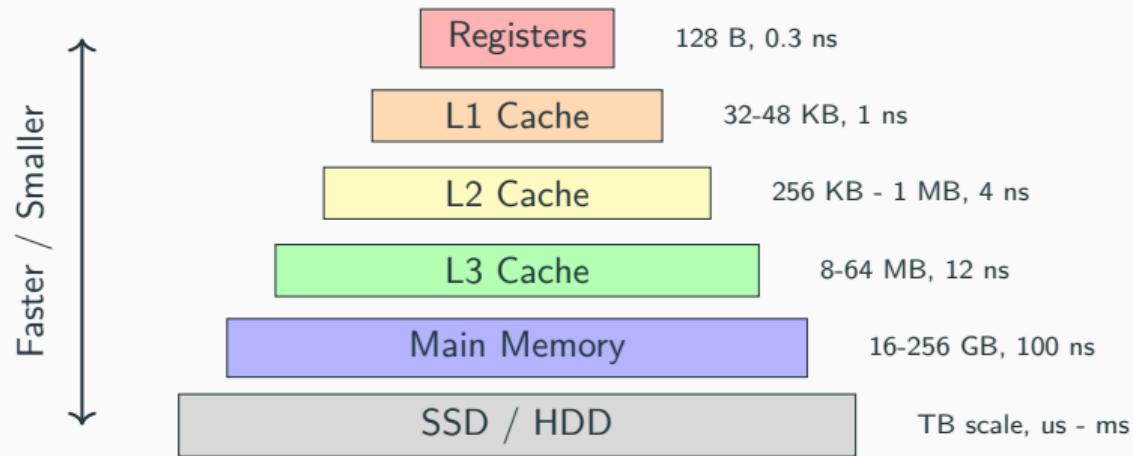
Where Does Data Live?

Data has its own journey through the memory hierarchy:



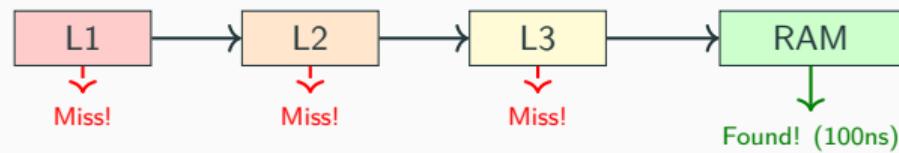
Each level is: **Smaller, Faster, More expensive**

The Memory Hierarchy



What Happens on a Memory Access?

Consider: `mov rdx, QWORD PTR [rbp-8]`



Cache Lines: The Unit of Transfer

You request array[0] (8 bytes), memory loads 64 bytes:

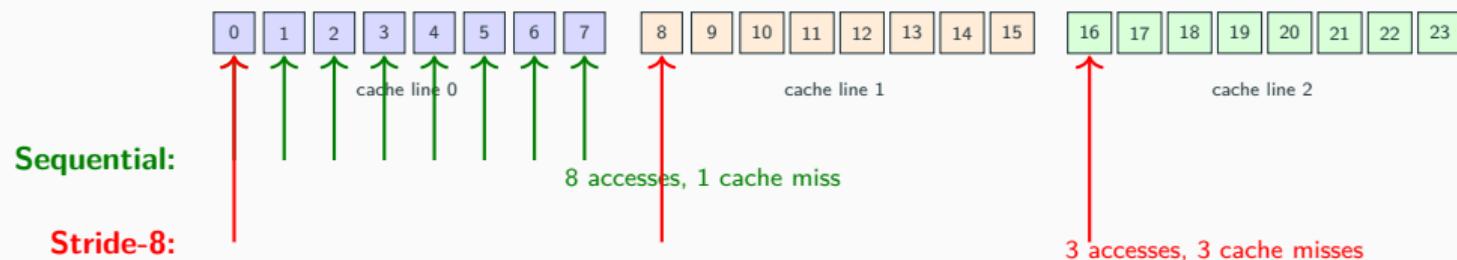


If you access array[1] next: **already in cache!**

Spatial locality: Sequential access is “free”.

Why Locality Matters

Memory: 24 elements across 3 cache lines



Sequential: All 8 accesses hit the same cache line → 1 miss total

Stride-8: Each access hits a different cache line → 1 miss per access ($8 \times$ worse!)

Two Types of Locality

Temporal Locality:

- Data accessed recently will likely be accessed again
- Example: loop counter, frequently-used variables
- Keep hot data in cache

Spatial Locality:

- Data near recently accessed data will likely be accessed
- Example: array traversal, struct fields
- Access data sequentially when possible

Data systems design is all about locality. Databases, ML systems, and every high-performance system we study in this course optimizes for these two properties. We will see this pattern again and again.

What Can Go Wrong?

Let's see locality in action with a simple experiment.

```
int matrix[1000][1000];
long sum = 0;

for (int i = 0; i < 1000; i++)
    for (int j = 0; j < 1000; j++)
        sum += matrix[i][j];

$ cd examples && make matrix_row && ./matrix_row
Sum: ..., Time: 2 ms
```

Try this one

```
int matrix[1000][1000];
long sum = 0;

for (int j = 0; j < 1000; j++)
    for (int i = 0; i < 1000; i++)
        sum += matrix[i][j];

$ cd examples && make matrix_col && ./matrix_col
Sum: ..., Time: 15 ms
```

Why 7× Slower?

Row-major traversal

```
for (i = 0; i < 1000; i++)  
    for (j = 0; j < 1000; j++)  
        sum += matrix[i][j];
```

Access: [0][0], [0][1], [0][2], ...

Linearized: 0, 1, 2, 3, 4, ...

Stride = 4 bytes

Column-major traversal

```
for (j = 0; j < 1000; j++)  
    for (i = 0; i < 1000; i++)  
        sum += matrix[i][j];
```

Access: [0][0], [1][0], [2][0], ...

Linearized: 0, 1000, 2000, 3000, ...

Stride = 4000 bytes

Why 7× Slower? (Cache View)

Row-major: stride = 4 bytes

Cache line = 64 bytes = 16 ints

Access 0, 1, 2, ..., 15 → **1 miss, 15 hits**

Then 16, 17, ..., 31 → **1 miss, 15 hits**

Miss rate: **6.25%**

Column-major: stride = 4000 bytes

Each access jumps 4000 bytes (62 cache lines!)

Access 0 → miss, load line

Access 1000 → miss, load line

Access 2000 → miss, load line

Miss rate: **100%**

The Result: 7× Slower

Traversal	Time	Cache behavior
Row-major	~2 ms	1 miss per 16 accesses
Column-major	~15 ms	1 miss per access

Same computation. Same data. 7× performance difference.

Measuring Cache Performance

Performance Counters (PMU)

Real-world applications are far more complex than our matrix example. Reasoning alone isn't enough—we need to **benchmark, profile, and measure** to understand where time is spent.

Modern CPUs have **Performance Monitoring Units** that count hardware events:

- Cache hits/misses at each level (L1, L2, L3)
- Branch predictions/mispredictions
- Instructions retired, cycles, etc.

On Linux, use `perf` to read these counters:

```
$ cd examples && make perf      # Run all with perf profiling
```

Measuring Row-Major Access

```
$ cd examples  
$ perf stat -e L1-dcache-loads,L1-dcache-load-misses ./matrix_row  
  
1,000,000,000      L1-dcache-loads  
62,500,000        L1-dcache-load-misses # 6.25% miss rate
```

Matches our prediction! 1M accesses, 1 miss per 16 accesses = 62.5K misses per iteration.

$$(1000 \text{ iterations} \times 62,500 = 62.5\text{M misses})$$

Measuring Column-Major Access

```
$ cd examples  
$ perf stat -e L1-dcache-loads,L1-dcache-load-misses ./matrix_col  
1,000,000,000      L1-dcache-loads  
1,000,000,000      L1-dcache-load-misses # 100% miss rate
```

Every single access misses L1 cache!

The data we need is never in cache because we jump 4000 bytes each time.

Cache Miss Summary

Traversal	L1 Loads	L1 Misses	Miss Rate
Row-major	1B	62.5M	6.25%
Column-major	1B	1B	100%

16× more L1 cache misses → 7× slower

(Not 16× slower because L2/L3 partially hide the latency)

Takeaway: perf lets you verify your mental model. When optimizing, measure first!

```
$ cd examples && make perf # Try it yourself!
```

Hazards: When Things Go Wrong

Pipeline Hazards

The pipeline assumes instructions flow smoothly. Three things break this:

1. **Data hazards:** Instruction needs result from previous instruction
2. **Control hazards:** Branch—which instruction comes next?
3. **Structural hazards:** Two instructions need same hardware

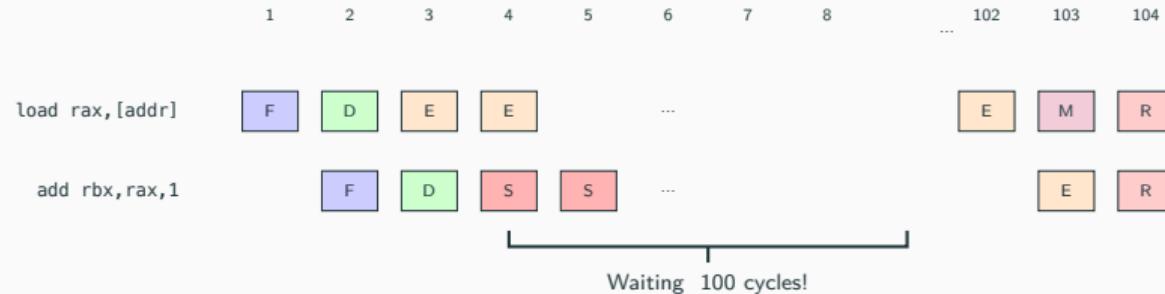
Pipeline Hazards

The pipeline assumes instructions flow smoothly. Three things break this:

1. **Data hazards:** Instruction needs result from previous instruction
2. **Control hazards:** Branch—which instruction comes next?
3. **Structural hazards:** Two instructions need same hardware

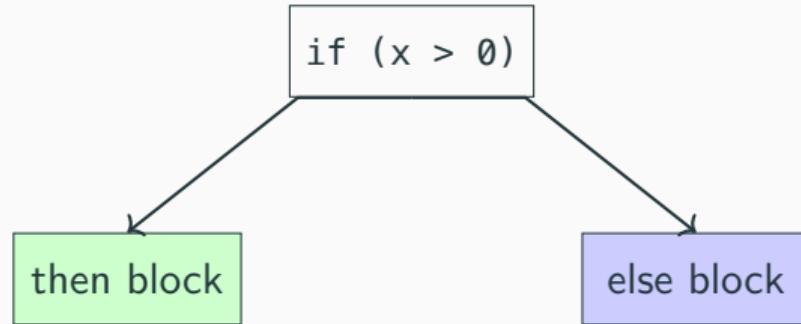
All of these are fundamentally about waiting for data.

Data Hazard Example



The add is **starved**—cannot execute until data arrives.

Control Hazard: Branches



Which path to fetch???

Pipeline must guess **before** condition is known!

Branch Misprediction Cost



Modern predictors are >95% accurate, **but** data-dependent branches can be unpredictable.

Which is Fastest? (Random Data)

Version A

```
for (i = 0; i < n; i++)
    if (data[i] > t)
        count++;
```

Version B

```
for (i = 0; i < n; i++)
    count += (data[i] > t);
```

Version C

```
sort(data, data + n);
for (i = 0; i < n; i++)
    if (data[i] > t)
        count++;
```

```
$ cd examples && make branch_test && ./branch_test
```

Version	Time
A (branching)	56 ms
B (branchless)	8 ms
C (sorted)	12 ms

Branchless is **7x faster** than branching on random data!

Data-Dependent Branches: The Problem

C++ code

```
for (int i = 0; i < n; i++) {  
    if (data[i] > threshold)  
        count++;  
}
```

On random data: branch taken ~50% of the time, **unpredictably**.

Branch predictor accuracy: ~50%
~15-20 cycle penalty per mispredict!

Assembly (x86-64)

```
loop:  
    mov eax, [rdi + rcx*4]  
    cmp eax, esi  
    jle skip ; BRANCH!  
    inc edx  
  
skip:  
    inc rcx  
    cmp rcx, r8  
    jl loop
```

The `jle skip` is the problem—CPU must guess which way to go.

Solution 1: Branchless Code

C++ code

```
for (int i = 0; i < n; i++) {  
    count += (data[i] > threshold);  
}
```

The comparison (`data[i] > threshold`) evaluates to 0 or 1.

No branch = no misprediction penalty!

Assembly (x86-64)

loop:

```
    mov  eax, [rdi + rcx*4]  
    cmp  eax, esi  
    setg al           ; al = 1 if greater  
    movzx eax, al  
    add  edx, eax     ; No branch!  
    inc  rcx  
    cmp  rcx, r8  
    jl   loop
```

setg + add replaces the conditional branch.

Solution 2: Sort First

C++ code

```
sort(data, data + n);

for (int i = 0; i < n; i++) {
    if (data[i] > threshold)
        count++;
}
```

After sorting: all values below threshold come first, then all above.

Branch pattern: FFFFF...TTTTT

Predictor accuracy: >99%

Why it works

Sorted data:

[1, 3, 5, 8, 12, 15, 20, ...]

^ threshold = 10

First iterations: always skip (predict: skip)

Later iterations: always take (predict: take)

Only ~1 misprediction at the transition!

Branch Performance Comparison

```
$ cd examples && perf stat -e branches,branch-misses ./branch_test
```

Approach	Time (random data)	Why
Branching	~15 ms	50% misprediction rate
Branchless	~3 ms	No branches to mispredict
Sort + Branch	~4 ms	Predictable pattern

Note: Sorting has $O(n \log n)$ cost, only worth it if you traverse multiple times or need sorted data anyway.

The Fundamental Problem

Every hazard is a data movement problem

- **Data hazard:** Waiting for data from memory
- **Control hazard:** Waiting for branch condition (data!)
- **Structural hazard:** Waiting for hardware to move data

The CPU can execute billions of operations per second...

...but only if the data is **in the right place at the right time**.

Strategies to Minimize Data Movement

1. Improve locality

- Sequential access patterns
- Keep working set small (fit in cache)
- Co-locate related data

2. Hide latency

- Prefetching (hardware and software)
- Out-of-order execution (do other work while waiting)
- Parallelism (multiple outstanding requests)

3. Reduce data volume

- Compression
- Filtering early (don't move data you'll discard)

Summary

Key Takeaways

1. **Instructions** flow through: Fetch → Decode → Execute → Retire
 - Pipelining and superscalar execution enable high throughput
 - But only if data is ready!
2. **Data** flows through: Disk → RAM → L3 → L2 → L1 → Registers
 - Each level faster but smaller
 - Cache lines (64 bytes) are the unit of transfer
3. **Hazards** occur when data isn't where it needs to be
 - Data hazards, control hazards—all about waiting for data
4. **Optimization is fundamentally about data movement**