

Lecture 8. SIMD and Vectorization

Introductions to Data Systems and Data Design

Ce Zhang

What is SIMD?

Flynn's Taxonomy

	Single Instruction	Multiple Instruction
Single Data	SISD (Uniprocessor)	MISD
Multiple Data	SIMD (Vector)	MIMD (Multicore)

Flynn's Taxonomy

	Single Instruction	Multiple Instruction
Single Data	SISD (Uniprocessor)	MISD
Multiple Data	SIMD (Vector)	MIMD (Multicore)

SIMD = Single Instruction, Multiple Data

One instruction operates on multiple data elements simultaneously.

SIMD: The Big Picture

Scalar:

$$a[0] + b[0] = c[0] \quad 1 \text{ operation}$$

SIMD (4-way):

$$\begin{array}{c} a[0] \\ a[1] \\ a[2] \\ a[3] \end{array} + \begin{array}{c} b[0] \\ b[1] \\ b[2] \\ b[3] \end{array} = \begin{array}{c} c[0] \\ c[1] \\ c[2] \\ c[3] \end{array} \quad \begin{array}{l} 4 \text{ operations} \\ \text{1 instruction!} \end{array}$$

Why SIMD Matters

Theoretical speedup: 4-8x for vectorizable code

Extension	Register Width	Doubles	Floats
SSE/SSE2	128 bits	2-way	4-way
AVX/AVX2	256 bits	4-way	8-way
AVX-512	512 bits	8-way	16-way

Why SIMD Matters

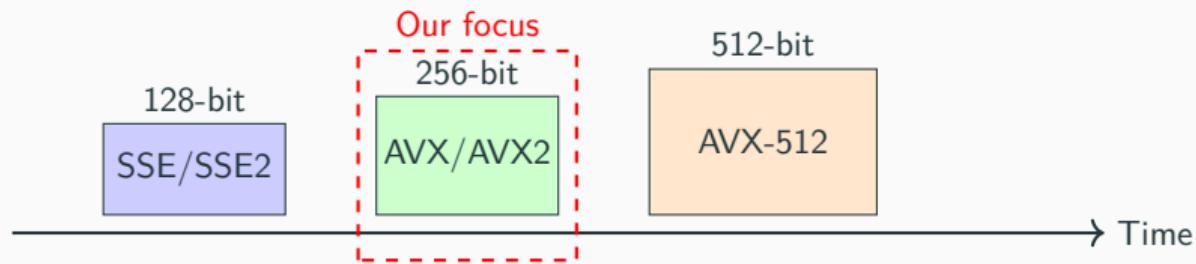
Theoretical speedup: 4-8x for vectorizable code

Extension	Register Width	Doubles	Floats
SSE/SSE2	128 bits	2-way	4-way
AVX/AVX2	256 bits	4-way	8-way
AVX-512	512 bits	8-way	16-way

SIMD + ILP together:

- ILP: Multiple *different* instructions in parallel
- SIMD: One instruction on *multiple data*
- Combine both for maximum throughput!

Evolution of x86 SIMD



AVX2 is widely available (Sandy Bridge 2011+). We'll focus on it.

AVX Registers and Data Types

AVX Vector Registers

16 registers: ymm0 - ymm15, each 256 bits wide

ymm0	d0	d1	d2	d3	4 doubles (64-bit each)
------	----	----	----	----	-------------------------

ymm1	f0	f1	f2	f3	f4	f5	f6	f7	8 floats (32-bit each)
------	----	----	----	----	----	----	----	----	------------------------

Same register, different interpretations!

Data Types in C/C++

```
#include <immmintrin.h> // AVX intrinsics header

__m256d d; // 4 doubles (256 bits)
__m256 f; // 8 floats (256 bits)
__m256i i; // 32 bytes / 16 shorts / 8 ints / 4 longs
```

Data Types in C/C++

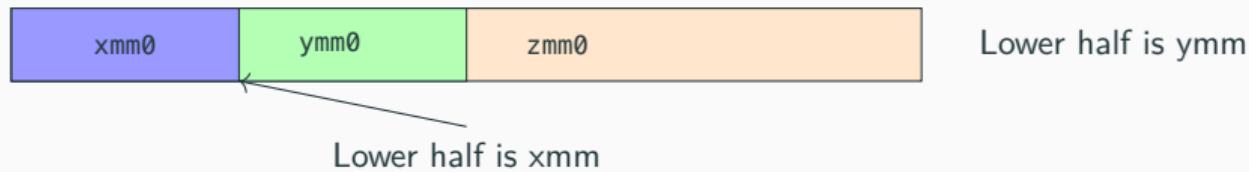
```
#include <immintrin.h> // AVX intrinsics header

__m256d d; // 4 doubles (256 bits)
__m256 f; // 8 floats (256 bits)
__m256i i; // 32 bytes / 16 shorts / 8 ints / 4 longs
```

Naming convention:

- `__m256` = 256-bit register
- `d` suffix = double precision
- `i` suffix = integer
- No suffix = single precision (float)

Register Hierarchy



`xmm0` is the lower 128 bits of `ymm0`, which is the lower 256 bits of `zmm0`.

AVX Intrinsics

What Are Intrinsics?

Intrinsics = C functions that map directly to assembly instructions

```
// C code with intrinsics
__m256d a = _mm256_load_pd(ptr);      // Load 4 doubles
__m256d b = _mm256_load_pd(ptr2);
__m256d c = _mm256_add_pd(a, b);      // Add 4 doubles
_mm256_store_pd(result, c);           // Store 4 doubles

; Generated assembly
vmovapd ymm0, [rdi]      ; Load aligned
vmovapd ymm1, [rsi]
vaddpd  ymm0, ymm0, ymm1 ; Packed double add
vmovapd [rdx], ymm0       ; Store aligned
```

Like writing assembly, but in C!

Intrinsic Naming Convention

mm256<operation>_<suffix>

Part	Meaning
_mm256	256-bit AVX operation
_mm	128-bit SSE operation
_mm512	512-bit AVX-512 operation

Suffix	Data Type
pd	Packed double (4 doubles)
ps	Packed single (8 floats)
epi32	Packed 32-bit integers
si256	256-bit integer

Load and Store

```
// Aligned load/store (pointer must be 32-byte aligned!)
__m256d a = _mm256_load_pd(ptr);      // Load 4 doubles
_mm256_store_pd(ptr, a);               // Store 4 doubles

// Unaligned load/store (any pointer)
__m256d b = _mm256_loadu_pd(ptr);     // Slower on some CPUs
_mm256_storeu_pd(ptr, b);
```

Load and Store

```
// Aligned load/store (pointer must be 32-byte aligned)
__m256d a = _mm256_load_pd(ptr);      // Load 4 doubles
_mm256_store_pd(ptr, a);              // Store 4 doubles

// Unaligned load/store (any pointer)
__m256d b = _mm256_loadu_pd(ptr);     // Slower on some CPUs
_mm256_storeu_pd(ptr, b);
```

Alignment matters!

- Aligned: pointer address divisible by 32
- Unaligned load on aligned boundary: seg fault!
- Modern CPUs: unaligned is nearly as fast (but prefer aligned)

Setting Constants

```
// Set all elements to the same value
__m256d ones = _mm256_set1_pd(1.0);
// Result: [1.0, 1.0, 1.0, 1.0]

// Set each element individually (note: reverse order!)
__m256d v = _mm256_set_pd(4.0, 3.0, 2.0, 1.0);
// Result: [1.0, 2.0, 3.0, 4.0] (LSB first!)

// Set to zero
__m256d zero = _mm256_setzero_pd();
// Result: [0.0, 0.0, 0.0, 0.0]
```

Setting Constants

```
// Set all elements to the same value
__m256d ones = _mm256_set1_pd(1.0);
// Result: [1.0, 1.0, 1.0, 1.0]

// Set each element individually (note: reverse order!)
__m256d v = _mm256_set_pd(4.0, 3.0, 2.0, 1.0);
// Result: [1.0, 2.0, 3.0, 4.0] (LSB first!)

// Set to zero
__m256d zero = _mm256_setzero_pd();
// Result: [0.0, 0.0, 0.0, 0.0]
```

Warning: `_mm256_set_pd` takes arguments in **reverse order!**

Arithmetic Operations

```
__m256d c;

c = _mm256_add_pd(a, b);    //  $c[i] = a[i] + b[i]$ 
c = _mm256_sub_pd(a, b);    //  $c[i] = a[i] - b[i]$ 
c = _mm256_mul_pd(a, b);    //  $c[i] = a[i] * b[i]$ 
c = _mm256_div_pd(a, b);    //  $c[i] = a[i] / b[i]$ 

c = _mm256_sqrt_pd(a);      //  $c[i] = \sqrt{a[i]}$ 
c = _mm256_max_pd(a, b);    //  $c[i] = \max(a[i], b[i])$ 
c = _mm256_min_pd(a, b);    //  $c[i] = \min(a[i], b[i])$ 

// Fused multiply-add (AVX2):  $c = a*b + c$ 
c = _mm256_fmadd_pd(a, b, c); // One instruction!
```

Performance: Latency and Throughput

Operation	Latency	Throughput
add_pd	4 cycles	2/cycle
mul_pd	4 cycles	2/cycle
div_pd	13-14 cycles	1/8 cycle
sqrt_pd	18 cycles	1/12 cycle
fmaadd_pd	4 cycles	2/cycle
load_pd	7 cycles	2/cycle

Performance: Latency and Throughput

Operation	Latency	Throughput
add_pd	4 cycles	2/cycle
mul_pd	4 cycles	2/cycle
div_pd	13-14 cycles	1/8 cycle
sqrt_pd	18 cycles	1/12 cycle
fmaadd_pd	4 cycles	2/cycle
load_pd	7 cycles	2/cycle

FMA is crucial: $a*b + c$ in one instruction with same latency as add alone!

Example: Sum Reduction (Revisited)

The Challenge

From ILP lecture: reductions have dependency chains.

```
double sum = 0;  
for (int i = 0; i < n; i++) {  
    sum += a[i]; // Each add depends on previous!  
}
```

SIMD alone doesn't help — we still have a single accumulator.

The Challenge

From ILP lecture: reductions have dependency chains.

```
double sum = 0;  
for (int i = 0; i < n; i++) {  
    sum += a[i]; // Each add depends on previous!  
}
```

SIMD alone doesn't help — we still have a single accumulator.

Solution: SIMD + multiple accumulators!

SIMD Sum: Step 1 - Vector Accumulator

```
__m256d vsum = _mm256_setzero_pd(); // [0, 0, 0, 0]

for (int i = 0; i < n; i += 4) {
    __m256d v = _mm256_loadu_pd(a + i);
    vsum = _mm256_add_pd(vsum, v); // 4 parallel sums!
}

// vsum = [sum0, sum1, sum2, sum3]
// Final: sum = sum0 + sum1 + sum2 + sum3
```

SIMD Sum: Step 1 - Vector Accumulator

```
__m256d vsum = _mm256_setzero_pd(); // [0, 0, 0, 0]

for (int i = 0; i < n; i += 4) {
    __m256d v = _mm256_loadu_pd(a + i);
    vsum = _mm256_add_pd(vsum, v); // 4 parallel sums!
}

// vsum = [sum0, sum1, sum2, sum3]
// Final: sum = sum0 + sum1 + sum2 + sum3
```

Now we have **4 independent accumulator chains** — 4x ILP!

SIMD Sum: Step 2 - Horizontal Reduction

```
// Reduce [sum0, sum1, sum2, sum3] to a single value
// Method 1: Extract and add scalars
double result[4];
_mm256_storeu_pd(result, vsum);
double sum = result[0] + result[1] + result[2] + result[3];
```

SIMD Sum: Step 2 - Horizontal Reduction

```
// Reduce [sum0, sum1, sum2, sum3] to a single value
// Method 1: Extract and add scalars
double result[4];
_mm256_storeu_pd(result, vsum);
double sum = result[0] + result[1] + result[2] + result[3];

// Method 2: Use horizontal add (faster)
__m128d low  = _mm256_castpd256_pd128(vsum);      // [sum0, sum1]
__m128d high = _mm256_extractf128_pd(vsum, 1);    // [sum2, sum3]
__m128d sum128 = _mm_add_pd(low, high);           // [sum0+sum2, sum1+sum3]
sum128 = _mm_hadd_pd(sum128, sum128);             // [total, total]
double sum = _mm_cvtsd_f64(sum128);                // Extract scalar
```

Complete SIMD Sum

```
// sum_v1.cpp
double sum_simd(double* a, int n) {
    __m256d vsum = _mm256_setzero_pd();

    for (int i = 0; i < n; i += 4) {
        __m256d v = _mm256_loadu_pd(a + i);
        vsum = _mm256_add_pd(vsum, v);
    }

    // Horizontal reduction
    __m128d low = _mm256_castpd256_pd128(vsum);
    __m128d high = _mm256_extractf128_pd(vsum, 1);
    __m128d sum128 = _mm_add_pd(low, high);
    sum128 = _mm_hadd_pd(sum128, sum128);
    return _mm_cvtsd_f64(sum128);
}
```

SIMD Sum: Even Better with Unrolling

```
// sum_v2.cpp
double sum_simd_unrolled(double* a, int n) {
    __m256d vsum1 = _mm256_setzero_pd();
    __m256d vsum2 = _mm256_setzero_pd(); // Two vector accumulators!

    for (int i = 0; i < n; i += 8) {
        __m256d v1 = _mm256_loadu_pd(a + i);
        __m256d v2 = _mm256_loadu_pd(a + i + 4);
        vsum1 = _mm256_add_pd(vsum1, v1);
        vsum2 = _mm256_add_pd(vsum2, v2);
    }

    __m256d vsum = _mm256_add_pd(vsum1, vsum2);
    // ... horizontal reduction ...
}
```

Sum Performance Comparison

```
$ cd examples && make sum
```

Version	Time
SIMD (1 accumulator)	104 ms
SIMD + 2x unrolled	52 ms

Sum Performance Comparison

```
$ cd examples && make sum
```

Version	Time
SIMD (1 accumulator)	104 ms
SIMD + 2x unrolled	52 ms

For maximum performance: **SIMD + unrolling** together!

Example: Dot Product

Scalar Dot Product

```
// dot_v1.cpp
double dot_scalar(double* a, double* b, int n) {
    double sum = 0;
    for (int i = 0; i < n; i++) {
        sum += a[i] * b[i];
    }
    return sum;
}
```

SIMD Dot Product with FMA

```
// dot_v2.cpp
double dot_simd(double* a, double* b, int n) {
    __m256d vsum = _mm256_setzero_pd();

    for (int i = 0; i < n; i += 4) {
        __m256d va = _mm256_loadu_pd(a + i);
        __m256d vb = _mm256_loadu_pd(b + i);
        // FMA: vsum = va * vb + vsum (one instruction!)
        vsum = _mm256_fmadd_pd(va, vb, vsum);
    }

    // Horizontal reduction (same as sum)
    __m128d low = _mm256_castpd256_pd128(vsum);
    __m128d high = _mm256_extractf128_pd(vsum, 1);
    __m128d sum128 = _mm_add_pd(low, high);
    sum128 = _mm_hadd_pd(sum128, sum128);
    return _mm_cvtsd_f64(sum128);
}
```

Dot Product Performance

```
$ cd examples && make dot
```

- FMA does multiply+add in one instruction
- 4 parallel multiply-adds per iteration
- Still need unrolling for maximum throughput

Handling Conditionals

The Problem with Branches

```
// Scalar version with branch
for (int i = 0; i < n; i++) {
    if (a[i] > 0.5)
        b[i] = a[i] + 1.0;
    else
        b[i] = a[i] - 1.0;
}
```

SIMD processes 4 elements at once — what if some need +1 and others need -1?

The Problem with Branches

```
// Scalar version with branch
for (int i = 0; i < n; i++) {
    if (a[i] > 0.5)
        b[i] = a[i] + 1.0;
    else
        b[i] = a[i] - 1.0;
}
```

SIMD processes 4 elements at once — what if some need +1 and others need -1?

Solution: Compute **both** paths, use **mask** to select results.

SIMD Comparison and Masking

```
$ cd examples && make cond
```

```
// conditional.cpp
__m256d threshold = _mm256_set1_pd(0.5);
__m256d ones = _mm256_set1_pd(1.0);
__m256d mones = _mm256_set1_pd(-1.0);
for (int i = 0; i < n; i += 4) {
    __m256d v = _mm256_loadu_pd(a + i);

    // Compare: creates mask (all 1s or all 0s per element)
    __m256d mask = _mm256_cmp_pd(v, threshold, _CMP_GT_OQ);

    // Blend: select from ones or mones based on mask
    __m256d delta = _mm256_blendv_pd(mones, ones, mask);

    __m256d result = _mm256_add_pd(v, delta);
    _mm256_storeu_pd(b + i, result);
}
```

Visualizing Masking

v:	0.3	0.8	0.2	0.9	
mask:	0x0	0xff..	0x0	0xff..	$(v > 0.5?)$
delta:	-1	+1	-1	+1	(blendv)
result:	-0.7	1.8	-0.8	1.9	

Compiler Vectorization

Can the Compiler Do This Automatically?

```
void add(double* a, double* b, double* c, int n) {  
    for (int i = 0; i < n; i++) {  
        c[i] = a[i] + b[i];  
    }  
}  
  
$ g++ -O3 -march=native add.cpp -S  
$ grep vaddpd add.s  
    vaddpd (%rsi,%rax), %ymm0, %ymm0    # Yes! Vectorized!
```

Can the Compiler Do This Automatically?

```
void add(double* a, double* b, double* c, int n) {  
    for (int i = 0; i < n; i++) {  
        c[i] = a[i] + b[i];  
    }  
}  
  
$ g++ -O3 -march=native add.cpp -S  
$ grep vaddpd add.s  
    vaddpd (%rsi,%rax), %ymm0, %ymm0    # Yes! Vectorized!
```

Yes, often! With `-O3 -march=native`, compilers auto-vectorize simple loops.

When to Use Intrinsics vs Compiler

Use compiler vectorization when:

- Simple loops (add, multiply arrays)
- Performance is “good enough”
- Portability matters

Use intrinsics when:

- Complex algorithms (reductions, shuffles)
- Need guaranteed vectorization
- Maximum performance required
- Compiler fails to vectorize

Practical Guidelines

SIMD Best Practices

1. **Use aligned memory** when possible (32-byte for AVX)
2. **Combine SIMD with unrolling** for maximum ILP
3. **Minimize shuffles** — they're expensive
4. **Process data in chunks of vector width** (4 doubles for AVX)
5. **Handle remainders** — what if $n \% 4 \neq 0$?
6. **Benchmark!** — SIMD doesn't always help (memory-bound code)

When SIMD Helps Most

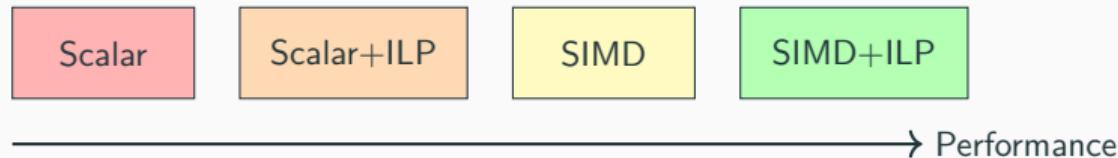
Good candidates:

- Element-wise operations (add, multiply arrays)
- Reductions (sum, max, dot product)
- Stencil computations
- Image processing
- Small matrix operations

Poor candidates:

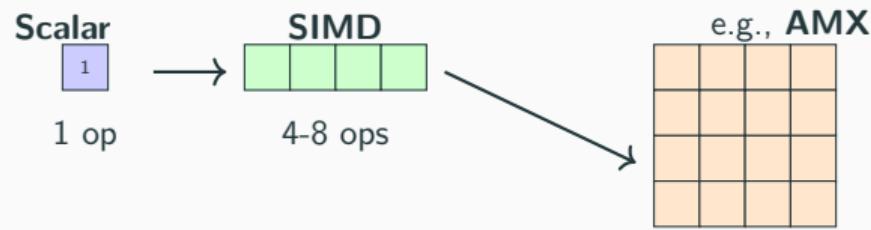
- Already memory-bound code
- Highly irregular access patterns
- Lots of branches/conditionals
- Gather/scatter intensive

Performance Hierarchy



Beyond SIMD: Scalar → Vector → Matrix

The Progression: Scalar → Vector → Matrix



$$16 \times 16 = 256 \text{ ops}$$

Apple AMX: Matrix Coprocessor

AMX = Apple Matrix eXtensions (M1/M2/M3 chips, undocumented but it is there)

- Dedicated matrix coprocessor (not just wider SIMD)
- Computes **16×16 outer products** in one instruction
- Used by Accelerate framework, PyTorch, etc.

Apple AMX: Matrix Coprocessor

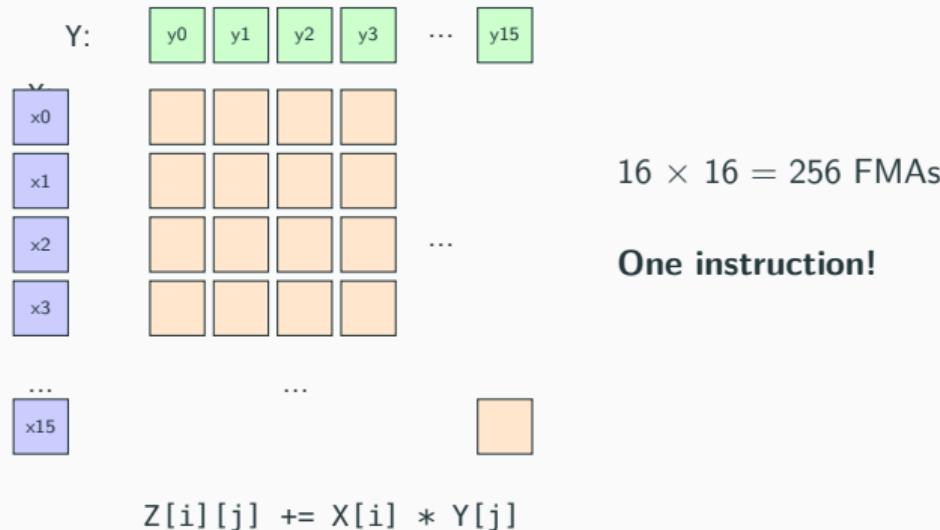
AMX = Apple Matrix eXtensions (M1/M2/M3 chips, undocumented but it is there)

- Dedicated matrix coprocessor (not just wider SIMD)
- Computes **16×16 outer products** in one instruction
- Used by Accelerate framework, PyTorch, etc.

Register layout:

Register	Size	Contents (fp32)
X[0-7]	64 bytes each	16 floats per register
Y[0-7]	64 bytes each	16 floats per register
Z	4 KB total	16×16 accumulator matrix

AMX Operation: Outer Product



AMX in C++: A Glimpse

```
#include "amx.h" // Apple's undocumented AMX "intrinsics"

// Compute 16x16 outer product: Z += X * Y^T
void amx_outer_product(float* X, float* Y, float* Z) {
    AMX_SET(); // Enable AMX coprocessor

    // Load 16 floats into X register 0
    AMX_LDX(ldx_operand(X, 0));

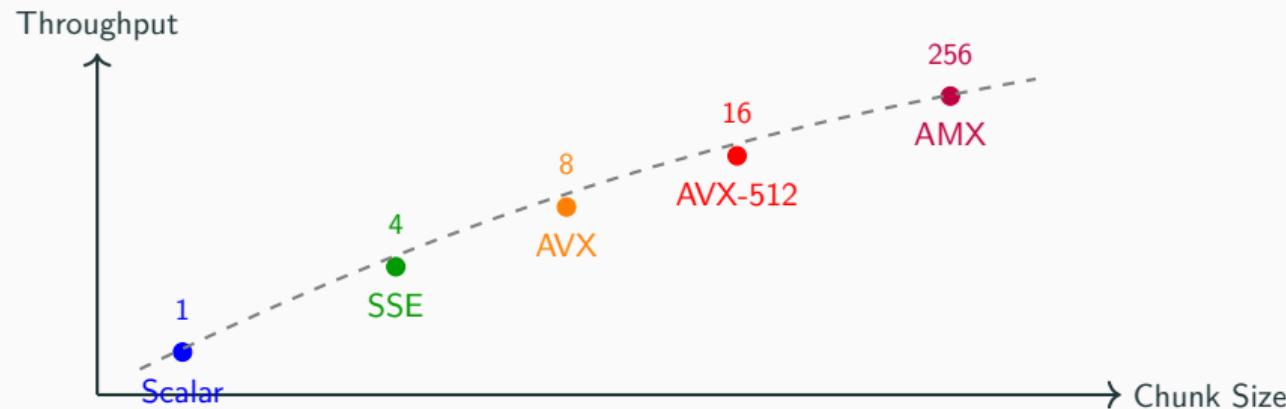
    // Load 16 floats into Y register 0
    AMX LDY(ldy_operand(Y, 0));

    // FMA: Z[i][j] += X[i] * Y[j] for all i,j in [0,16]
    AMX_FMA32(fma32_operand(0, 0, 0));

    // Store result (64 bytes per row, 16 rows)
    for (int row = 0; row < 16; row++)
        AMX_STZ(stz_operand(Z + row * 16, row));

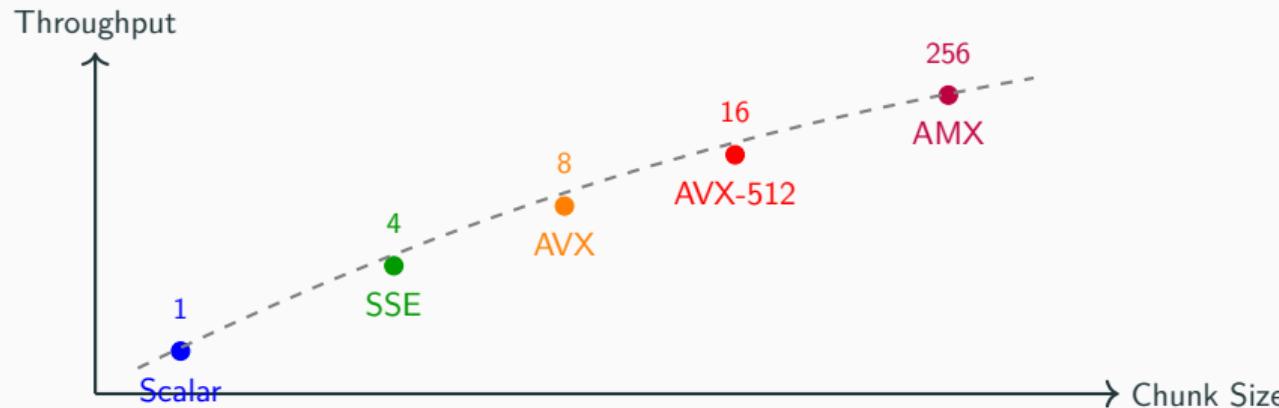
    AMX_CLR(); // Disable AMX
}
```

The Big Picture: Processing in Chunks



Trend: Increasing parallelism through larger “chunks” of computation

The Big Picture: Processing in Chunks



Trend: Increasing parallelism through larger “chunks” of computation

On GPUs: This becomes **Tensor Cores** (NVIDIA)

```
// PTX instruction: 16x8x16 matrix multiply = 4096 FLOPs!
mma.sync.aligned.m16n8k16.row.col.f16.f16.f16.f16
    {d0,d1,d2,d3}, {a0,a1,a2,a3}, {b0,b1}, {c0,c1,c2,c3};
```

Summary

Key Takeaways

1. **SIMD** = process 4-8 elements with one instruction
 - AVX: 4 doubles or 8 floats per instruction
2. **Intrinsics** give you direct control over SIMD
 - `_mm256_add_pd`, `_mm256_load_pd`, etc.
3. **SIMD + ILP** together for maximum performance
 - Use multiple vector accumulators
4. **Compiler can auto-vectorize** simple cases
 - But intrinsics give you more control
5. **Alignment** matters — use 32-byte aligned memory

Try It Yourself

```
# Run all examples
```

```
$ cd examples && make run
```

```
# Build with -O3 for comparison
```

```
$ cd examples && make fast && make run
```

```
# Use perf to measure
```

```
$ cd examples && make perf
```

Reference: Intel Intrinsics Guide