

Lecture 8. SIMD and Vectorization

Introductions to Data Systems and Data Design

Ce Zhang

What is SIMD?

Flynn's Taxonomy

	Single Instruction	Multiple Instruction
Single Data	SISD (Uniprocessor)	MISD
Multiple Data	SIMD (Vector)	MIMD (Multicore)

Flynn's Taxonomy

	Single Instruction	Multiple Instruction
Single Data	SISD (Uniprocessor)	MISD
Multiple Data	SIMD (Vector)	MIMD (Multicore)

SIMD = Single Instruction, Multiple Data

One instruction operates on multiple data elements simultaneously.

SIMD: The Big Picture

Scalar:

$$a[0] + b[0] = c[0] \quad 1 \text{ operation}$$

SIMD (4-way):

$$\begin{array}{c} a[0] \\ a[1] \\ a[2] \\ a[3] \end{array} + \begin{array}{c} b[0] \\ b[1] \\ b[2] \\ b[3] \end{array} = \begin{array}{c} c[0] \\ c[1] \\ c[2] \\ c[3] \end{array} \quad \begin{array}{l} 4 \text{ operations} \\ \textbf{1 instruction!} \end{array}$$

Why SIMD Matters

Theoretical speedup: 4-8x for vectorizable code

Extension	Register Width	Doubles	Floats
SSE/SSE2	128 bits	2-way	4-way
AVX/AVX2	256 bits	4-way	8-way
AVX-512	512 bits	8-way	16-way

Why SIMD Matters


Theoretical speedup: 4-8x for vectorizable code

Extension	Register Width	Doubles	Floats
SSE/SSE2	128 bits	2-way	4-way
AVX/AVX2	256 bits	4-way	8-way
AVX-512	512 bits	8-way	16-way

SIMD + ILP together:

- ILP: Multiple *different* instructions in parallel
- SIMD: One instruction on *multiple data*
- Combine both for maximum throughput!

Evolution of x86 SIMD



AVX2 is widely available (Sandy Bridge 2011+). We'll focus on it.

AVX Registers and Data Types

AVX Vector Registers

16 registers: ymm0 - ymm15, each 256 bits wide

ymm0	d0	d1	d2	d3	4 doubles (64-bit each)
------	----	----	----	----	-------------------------

ymm1	f0	f1	f2	f3	f4	f5	f6	f7	8 floats (32-bit each)
------	----	----	----	----	----	----	----	----	------------------------

Same register, different interpretations!

Data Types in C/C++

```
#include <immmintrin.h> // AVX intrinsics header

__m256d d; // 4 doubles (256 bits)
__m256 f; // 8 floats (256 bits)
__m256i i; // 32 bytes / 16 shorts / 8 ints / 4 longs
```

Data Types in C/C++


```
#include <immintrin.h> // AVX intrinsics header

__m256d d; // 4 doubles (256 bits)
__m256 f; // 8 floats (256 bits)
__m256i i; // 32 bytes / 16 shorts / 8 ints / 4 longs
```

Naming convention:

- `__m256` = 256-bit register
- `d` suffix = double precision
- `i` suffix = integer
- No suffix = single precision (float)

Register Hierarchy



xmm0 is the lower 128 bits of ymm0, which is the lower 256 bits of zmm0.

AVX Intrinsics

What Are Intrinsics?

Intrinsics = C functions that map directly to assembly instructions

```
// C code with intrinsics
__m256d a = _mm256_load_pd(ptr);      // Load 4 doubles
__m256d b = _mm256_load_pd(ptr2);
__m256d c = _mm256_add_pd(a, b);      // Add 4 doubles
_mm256_store_pd(result, c);           // Store 4 doubles

; Generated assembly
vmovapd ymm0, [rdi]      ; Load aligned
vmovapd ymm1, [rsi]
vaddpd  ymm0, ymm0, ymm1 ; Packed double add
vmovapd [rdx], ymm0       ; Store aligned
```

Like writing assembly, but in C!

Intrinsic Naming Convention

mm256<operation>_<suffix>

Part	Meaning
_mm256	256-bit AVX operation
_mm	128-bit SSE operation
_mm512	512-bit AVX-512 operation

Suffix	Data Type
pd	Packed double (4 doubles)
ps	Packed single (8 floats)
epi32	Packed 32-bit integers
si256	256-bit integer

Load and Store

```
// Aligned load/store (pointer must be 32-byte aligned!)
__m256d a = _mm256_load_pd(ptr);      // Load 4 doubles
_mm256_store_pd(ptr, a);               // Store 4 doubles

// Unaligned load/store (any pointer)
__m256d b = _mm256_loadu_pd(ptr);     // Slower on some CPUs
_mm256_storeu_pd(ptr, b);
```

Load and Store

```
// Aligned load/store (pointer must be 32-byte aligned)
__m256d a = _mm256_load_pd(ptr);      // Load 4 doubles
_mm256_store_pd(ptr, a);               // Store 4 doubles

// Unaligned load/store (any pointer)
__m256d b = _mm256_loadu_pd(ptr);     // Slower on some CPUs
_mm256_storeu_pd(ptr, b);
```

Alignment matters!

- Aligned: pointer address divisible by 32
- Unaligned load on aligned boundary: seg fault!
- Modern CPUs: unaligned is nearly as fast (but prefer aligned)

Setting Constants

```
// Set all elements to the same value
__m256d ones = _mm256_set1_pd(1.0);
// Result: [1.0, 1.0, 1.0, 1.0]

// Set each element individually (note: reverse order!)
__m256d v = _mm256_set_pd(4.0, 3.0, 2.0, 1.0);
// Result: [1.0, 2.0, 3.0, 4.0] (LSB first!)

// Set to zero
__m256d zero = _mm256_setzero_pd();
// Result: [0.0, 0.0, 0.0, 0.0]
```

Setting Constants

```
// Set all elements to the same value
__m256d ones = _mm256_set1_pd(1.0);
// Result: [1.0, 1.0, 1.0, 1.0]

// Set each element individually (note: reverse order!)
__m256d v = _mm256_set_pd(4.0, 3.0, 2.0, 1.0);
// Result: [1.0, 2.0, 3.0, 4.0] (LSB first!)

// Set to zero
__m256d zero = _mm256_setzero_pd();
// Result: [0.0, 0.0, 0.0, 0.0]
```

Warning: `_mm256_set_pd` takes arguments in **reverse order!**

Arithmetic Operations

```
__m256d c;

c = _mm256_add_pd(a, b);    //  $c[i] = a[i] + b[i]$ 
c = _mm256_sub_pd(a, b);    //  $c[i] = a[i] - b[i]$ 
c = _mm256_mul_pd(a, b);    //  $c[i] = a[i] * b[i]$ 
c = _mm256_div_pd(a, b);    //  $c[i] = a[i] / b[i]$ 

c = _mm256_sqrt_pd(a);      //  $c[i] = \sqrt{a[i]}$ 
c = _mm256_max_pd(a, b);    //  $c[i] = \max(a[i], b[i])$ 
c = _mm256_min_pd(a, b);    //  $c[i] = \min(a[i], b[i])$ 

// Fused multiply-add (AVX2):  $c = a*b + c$ 
c = _mm256_fmadd_pd(a, b, c); // One instruction!
```

Performance: Latency and Throughput

Operation	Latency	Throughput
add_pd	4 cycles	2/cycle
mul_pd	4 cycles	2/cycle
div_pd	13-14 cycles	1/8 cycle
sqrt_pd	18 cycles	1/12 cycle
fmaadd_pd	4 cycles	2/cycle
load_pd	7 cycles	2/cycle

Performance: Latency and Throughput

Operation	Latency	Throughput
add_pd	4 cycles	2/cycle
mul_pd	4 cycles	2/cycle
div_pd	13-14 cycles	1/8 cycle
sqrt_pd	18 cycles	1/12 cycle
fmaadd_pd	4 cycles	2/cycle
load_pd	7 cycles	2/cycle

FMA is crucial: $a*b + c$ in one instruction with same latency as add alone!

Example: Vector Addition

Scalar Version

```
// add_scalar.cpp
void add_scalar(double* a, double* b, double* c, int n) {
    for (int i = 0; i < n; i++) {
        c[i] = a[i] + b[i];
    }
}

$ g++ -O1 add_scalar.cpp -o add_scalar && ./add_scalar
Time: 85 ms (100M elements)
```


SIMD Version

```
// add_simd.cpp
#include <immintrin.h>

void add_simd(double* a, double* b, double* c, int n) {
    for (int i = 0; i < n; i += 4) {
        __m256d va = _mm256_loadu_pd(a + i);
        __m256d vb = _mm256_loadu_pd(b + i);
        __m256d vc = _mm256_add_pd(va, vb);
        _mm256_storeu_pd(c + i, vc);
    }
}
```

```
$ g++ -O1 -mavx2 add_simd.cpp -o add_simd && ./add_simd
Time: 42 ms (100M elements)
```

Visualizing the SIMD Loop



Each loop iteration processes **4 elements** with one SIMD instruction.

Example: Sum Reduction (Revisited)

The Challenge

From ILP lecture: reductions have dependency chains.

```
double sum = 0;  
for (int i = 0; i < n; i++) {  
    sum += a[i]; // Each add depends on previous!  
}
```

SIMD alone doesn't help — we still have a single accumulator.

The Challenge

From ILP lecture: reductions have dependency chains.

```
double sum = 0;  
for (int i = 0; i < n; i++) {  
    sum += a[i]; // Each add depends on previous!  
}
```

SIMD alone doesn't help — we still have a single accumulator.

Solution: SIMD + multiple accumulators!

SIMD Sum: Step 1 - Vector Accumulator

```
__m256d vsum = _mm256_setzero_pd(); // [0, 0, 0, 0]

for (int i = 0; i < n; i += 4) {
    __m256d v = _mm256_loadu_pd(a + i);
    vsum = _mm256_add_pd(vsum, v); // 4 parallel sums!
}

// vsum = [sum0, sum1, sum2, sum3]
// Final: sum = sum0 + sum1 + sum2 + sum3
```

SIMD Sum: Step 1 - Vector Accumulator

```
__m256d vsum = _mm256_setzero_pd(); // [0, 0, 0, 0]

for (int i = 0; i < n; i += 4) {
    __m256d v = _mm256_loadu_pd(a + i);
    vsum = _mm256_add_pd(vsum, v); // 4 parallel sums!
}

// vsum = [sum0, sum1, sum2, sum3]
// Final: sum = sum0 + sum1 + sum2 + sum3
```

Now we have **4 independent accumulator chains** — 4x ILP!

SIMD Sum: Step 2 - Horizontal Reduction

```
// Reduce [sum0, sum1, sum2, sum3] to a single value
// Method 1: Extract and add scalars
double result[4];
_mm256_storeu_pd(result, vsum);
double sum = result[0] + result[1] + result[2] + result[3];
```

SIMD Sum: Step 2 - Horizontal Reduction

```
// Reduce [sum0, sum1, sum2, sum3] to a single value
// Method 1: Extract and add scalars
double result[4];
_mm256_storeu_pd(result, vsum);
double sum = result[0] + result[1] + result[2] + result[3];

// Method 2: Use horizontal add (faster)
__m128d low  = _mm256_castpd256_pd128(vsum);      // [sum0, sum1]
__m128d high = _mm256_extractf128_pd(vsum, 1);    // [sum2, sum3]
__m128d sum128 = _mm_add_pd(low, high);           // [sum0+sum2, sum1+sum3]
sum128 = _mm_hadd_pd(sum128, sum128);             // [total, total]
double sum = _mm_cvtsd_f64(sum128);                // Extract scalar
```

Complete SIMD Sum

```
double sum_simd(double* a, int n) {
    __m256d vsum = _mm256_setzero_pd();

    for (int i = 0; i < n; i += 4) {
        __m256d v = _mm256_loadu_pd(a + i);
        vsum = _mm256_add_pd(vsum, v);
    }

// Horizontal reduction
    __m128d low = _mm256_castpd256_pd128(vsum);
    __m128d high = _mm256_extractf128_pd(vsum, 1);
    __m128d sum128 = _mm_add_pd(low, high);
    sum128 = _mm_hadd_pd(sum128, sum128);
    return _mm_cvtsd_f64(sum128);
```

SIMD Sum: Even Better with Unrolling

```
double sum_simd_unrolled(double* a, int n) {
    __m256d vsum1 = _mm256_setzero_pd();
    __m256d vsum2 = _mm256_setzero_pd(); // Two vector accumulators!

    for (int i = 0; i < n; i += 8) {
        __m256d v1 = _mm256_loadu_pd(a + i);
        __m256d v2 = _mm256_loadu_pd(a + i + 4);
        vsum1 = _mm256_add_pd(vsum1, v1);
        vsum2 = _mm256_add_pd(vsum2, v2);
    }

    __m256d vsum = _mm256_add_pd(vsum1, vsum2);
    // ... horizontal reduction ...
}
```

Sum Performance Comparison

Version	Time	Speedup
Scalar naive	152 ms	1x
Scalar 4x unrolled	42 ms	3.6x
SIMD (1 accumulator)	42 ms	3.6x
SIMD + 2x unrolled	25 ms	6x

Sum Performance Comparison

Version	Time	Speedup
Scalar naive	152 ms	1x
Scalar 4x unrolled	42 ms	3.6x
SIMD (1 accumulator)	42 ms	3.6x
SIMD + 2x unrolled	25 ms	6x

Key insight: SIMD with one accumulator = scalar with 4 accumulators

For maximum performance: **SIMD + unrolling** together!

Example: Dot Product

Scalar Dot Product

```
double dot_scalar(double* a, double* b, int n) {  
    double sum = 0;  
    for (int i = 0; i < n; i++) {  
        sum += a[i] * b[i];  
    }  
    return sum;  
}
```

SIMD Dot Product with FMA

```
double dot_simd(double* a, double* b, int n) {
    __m256d vsum = _mm256_setzero_pd();

    for (int i = 0; i < n; i += 4) {
        __m256d va = _mm256_loadu_pd(a + i);
        __m256d vb = _mm256_loadu_pd(b + i);
        // FMA: vsum = va * vb + vsum (one instruction!)
        vsum = _mm256_fmadd_pd(va, vb, vsum);
    }

    // Horizontal reduction (same as sum)
    __m128d low = _mm256_castpd256_pd128(vsum);
    __m128d high = _mm256_extractf128_pd(vsum, 1);
    __m128d sum128 = _mm_add_pd(low, high);
```

Dot Product Performance

```
$ g++ -O1 -mavx2 -mfma dot.cpp -o dot && ./dot  
Scalar: 168 ms  
SIMD:    35 ms
```

4.8x speedup!

- FMA does multiply+add in one instruction
- 4 parallel multiply-adds per iteration
- Still need unrolling for maximum throughput

Handling Conditionals

The Problem with Branches

```
// Scalar version with branch
for (int i = 0; i < n; i++) {
    if (a[i] > 0.5)
        b[i] = a[i] + 1.0;
    else
        b[i] = a[i] - 1.0;
}
```

SIMD processes 4 elements at once — what if some need +1 and others need -1?

The Problem with Branches

```
// Scalar version with branch
for (int i = 0; i < n; i++) {
    if (a[i] > 0.5)
        b[i] = a[i] + 1.0;
    else
        b[i] = a[i] - 1.0;
}
```

SIMD processes 4 elements at once — what if some need +1 and others need -1?

Solution: Compute **both** paths, use **mask** to select results.

SIMD Comparison and Masking

```
_m256d threshold = _mm256_set1_pd(0.5);
_m256d ones = _mm256_set1_pd(1.0);
_m256d mones = _mm256_set1_pd(-1.0);

for (int i = 0; i < n; i += 4) {
    _m256d v = _mm256_loadu_pd(a + i);

    // Compare: creates mask (all 1s or all 0s per element)
    _m256d mask = _mm256_cmp_pd(v, threshold, _CMP_GT_0Q);

    // Blend: select from ones or mones based on mask
    _m256d delta = _mm256_blendv_pd(mones, ones, mask);

    _m256d result = _mm256_add_pd(v, delta);
```

Visualizing Masking

v:	0.3	0.8	0.2	0.9	
mask:	0x0	0xff..	0x0	0xff..	$(v > 0.5?)$
delta:	-1	+1	-1	+1	(blendv)
result:	-0.7	1.8	-0.8	1.9	

Compiler Vectorization

Can the Compiler Do This Automatically?

```
void add(double* a, double* b, double* c, int n) {  
    for (int i = 0; i < n; i++) {  
        c[i] = a[i] + b[i];  
    }  
}  
  
$ g++ -O3 -march=native add.cpp -S  
$ grep vaddpd add.s  
    vaddpd (%rsi,%rax), %ymm0, %ymm0    # Yes! Vectorized!
```

Can the Compiler Do This Automatically?

```
void add(double* a, double* b, double* c, int n) {  
    for (int i = 0; i < n; i++) {  
        c[i] = a[i] + b[i];  
    }  
}  
  
$ g++ -O3 -march=native add.cpp -S  
$ grep vaddpd add.s  
    vaddpd (%rsi,%rax), %ymm0, %ymm0    # Yes! Vectorized!
```

Yes, often! With `-O3 -march=native`, compilers auto-vectorize simple loops.

When Auto-Vectorization Fails

1. Aliasing — pointers might overlap

```
void add(double* a, double* b, double* c, int n) {  
    for (int i = 0; i < n; i++)  
        c[i] = a[i] + b[i]; // What if c overlaps a or b?  
}
```

Fix: Use restrict keyword or #pragma

```
void add(double* restrict a, double* restrict b,  
        double* restrict c, int n) { ... }
```

When Auto-Vectorization Fails (cont'd)

2. Complex control flow

```
while (a[i] > 0) { // Not a countable loop!
    ...
}
```

3. Function calls in loop

4. Floating-point associativity (for reductions)

```
sum += a[i]; // Can't reorder without -ffast-math
```

Fix: Use `-ffast-math` or `-fassociative-math`

Helping the Compiler

```
// Tell compiler there's no aliasing
#pragma GCC ivdep
for (int i = 0; i < n; i++)
    c[i] = a[i] + b[i];

// Force vectorization (GCC)
#pragma GCC optimize("tree-vectorize")

// Check if compiler vectorized
$ g++ -O3 -fopt-info-vec-optimized file.cpp
file.cpp:5: optimized: loop vectorized using 32 byte vectors
```

When to Use Intrinsics vs Compiler

Use compiler vectorization when:

- Simple loops (add, multiply arrays)
- Performance is “good enough”
- Portability matters

Use intrinsics when:

- Complex algorithms (reductions, shuffles)
- Need guaranteed vectorization
- Maximum performance required
- Compiler fails to vectorize

Alignment

Why Alignment Matters

32-byte alignment = address divisible by 32

```
// Unaligned (any address)
```

```
double* a = (double*)malloc(n * sizeof(double));
```

```
// Aligned to 32 bytes
```

```
double* a = (double*)aligned_alloc(32, n * sizeof(double));
```

Why Alignment Matters

32-byte alignment = address divisible by 32

```
// Unaligned (any address)
```

```
double* a = (double*)malloc(n * sizeof(double));
```

```
// Aligned to 32 bytes
```

```
double* a = (double*)aligned_alloc(32, n * sizeof(double));
```

Performance impact:

- Modern CPUs: unaligned nearly as fast
- But: `_mm256_load_pd` on unaligned = **crash!**
- Aligned loads enable more optimizations

Allocating Aligned Memory

```
// C11 standard
double* a = aligned_alloc(32, n * sizeof(double));
```

```
// POSIX
double* a;
posix_memalign((void**)&a, 32, n * sizeof(double));
```

```
// C++ 17
double* a = static_cast<double*>(
    std::aligned_alloc(32, n * sizeof(double)));
```

```
// On stack
alignas(32) double a[1024];
```

Safe Loading Pattern

```
void process(double* a, int n) {
    int i = 0;

    // Handle unaligned prefix
    while (((uintptr_t)(a + i) & 31) != 0 && i < n) {
        // Scalar processing
        process_one(a[i]);
        i++;
    }

    // Main loop: aligned
    for (; i + 4 <= n; i += 4) {
        __m256d v = _mm256_load_pd(a + i); // Safe!
        // ...
    }
}
```

Practical Guidelines

SIMD Best Practices

1. **Use aligned memory** when possible (32-byte for AVX)
2. **Combine SIMD with unrolling** for maximum ILP
3. **Minimize shuffles** — they're expensive
4. **Process data in chunks of vector width** (4 doubles for AVX)
5. **Handle remainders** — what if $n \% 4 \neq 0$?
6. **Benchmark!** — SIMD doesn't always help (memory-bound code)

When SIMD Helps Most

Good candidates:

- Element-wise operations (add, multiply arrays)
- Reductions (sum, max, dot product)
- Stencil computations
- Image processing
- Small matrix operations

Poor candidates:

- Already memory-bound code
- Highly irregular access patterns
- Lots of branches/conditionals
- Gather/scatter intensive

Complete Example: Array Operations

```
#include <immmintrin.h>
#include <cstdlib>

void saxpy_simd(float* y, float a, float* x, int n) {
    __m256 va = _mm256_set1_ps(a); // Broadcast a


    int i = 0;
    for (; i + 8 <= n; i += 8) {
        __m256 vx = _mm256_loadu_ps(x + i);
        __m256 vy = _mm256_loadu_ps(y + i);
        vy = _mm256_fmadd_ps(va, vx, vy); // y = a*x + y
        _mm256_storeu_ps(y + i, vy);
    }
}
```

Summary

Key Takeaways

1. **SIMD** = process 4-8 elements with one instruction
 - AVX: 4 doubles or 8 floats per instruction
2. **Intrinsics** give you direct control over SIMD
 - `_mm256_add_pd`, `_mm256_load_pd`, etc.
3. **SIMD + ILP** together for maximum performance
 - Use multiple vector accumulators
4. **Compiler can auto-vectorize** simple cases
 - But intrinsics give you more control
5. **Alignment** matters — use 32-byte aligned memory

Performance Hierarchy



Try It Yourself

Compile with AVX support

```
$ g++ -O3 -mavx2 -mfma program.cpp -o program
```

Check what instructions are generated

```
$ g++ -O3 -mavx2 -S program.cpp
```

```
$ grep -E "vadd|vmul|vfma" program.s
```

See if compiler auto-vectorized

```
$ g++ -O3 -mavx2 -fopt-info-vec program.cpp
```

Use perf to measure

```
$ perf stat ./program
```

Reference: Intel Intrinsics Guide