

# Lecture 10. The Roofline Model

Introductions to Data Systems and Data Design

---

Ce Zhang

# Introduction

---

# The Fundamental Question

When optimizing code, we need to know:

**Is my code limited by compute or by memory bandwidth?**

# The Fundamental Question

When optimizing code, we need to know:

**Is my code limited by compute or by memory bandwidth?**

- If **compute bound**: Optimize arithmetic (SIMD, ILP, better algorithms)
- If **memory bound**: Optimize data movement (blocking, prefetching, data layout)

# The Fundamental Question

When optimizing code, we need to know:


**Is my code limited by compute or by memory bandwidth?**

- If **compute bound**: Optimize arithmetic (SIMD, ILP, better algorithms)
- If **memory bound**: Optimize data movement (blocking, prefetching, data layout)

The **Roofline Model** answers this question visually and quantitatively.

# The Two Ceilings

Every program faces two fundamental limits:



## Platform Parameters

**Peak Performance  $\pi$**  [flops/cycle or GFLOPS]:

$$\pi = \text{cores} \times \text{SIMD width} \times \text{FMA units} \times \text{ops/FMA}$$

## Platform Parameters

**Peak Performance  $\pi$**  [flops/cycle or GFLOPS]:

$$\pi = \text{cores} \times \text{SIMD width} \times \text{FMA units} \times \text{ops/FMA}$$

**Example:** Intel Core i7 (Skylake), single core:

- 2 FMA units
- 4-way SIMD (AVX, doubles)
- 2 ops per FMA (multiply + add)

$$\pi = 1 \times 4 \times 2 \times 2 = 16 \text{ flops/cycle}$$

## Platform Parameters

**Peak Bandwidth**  $\beta$  [bytes/cycle or GB/s]:

Theoretical bandwidth is determined by hardware specs:

$$\beta_{\text{theo}} = \text{channels} \times \text{data rate} \times \text{bus width}$$

## Platform Parameters

**Peak Bandwidth**  $\beta$  [bytes/cycle or GB/s]:

Theoretical bandwidth is determined by hardware specs:

$$\beta_{\text{theo}} = \text{channels} \times \text{data rate} \times \text{bus width}$$

**Example:** DDR4-2400, dual channel:

- $2 \text{ channels} \times 2400 \text{ MT/s} \times 8 \text{ bytes} = 38.4 \text{ GB/s}$

## Algorithm Parameters

**Work**  $W(n)$  [flops]: Total floating-point operations

**Data Movement**  $Q(n)$  [bytes]: Bytes transferred between cache and memory

## Algorithm Parameters

**Work**  $W(n)$  [flops]: Total floating-point operations

**Data Movement**  $Q(n)$  [bytes]: Bytes transferred between cache and memory

**Operational Intensity:**

$$I(n) = \frac{W(n)}{Q(n)} \quad [\text{flops}/\text{byte}]$$

## Algorithm Parameters

**Work**  $W(n)$  [flops]: Total floating-point operations

**Data Movement**  $Q(n)$  [bytes]: Bytes transferred between cache and memory

**Operational Intensity:**

$$I(n) = \frac{W(n)}{Q(n)} \quad [\text{flops}/\text{byte}]$$

**Performance:**

$$P(n) = \frac{W(n)}{T(n)} \quad [\text{flops}/\text{cycle}]$$

where  $T(n)$  is runtime in cycles.

## Deriving the Roofline

---

# The Two Bounds

## Bound 1: Compute Bound

$$P \leq \pi$$

Performance cannot exceed peak compute throughput.

# The Two Bounds

## Bound 1: Compute Bound

$$P \leq \pi$$

Performance cannot exceed peak compute throughput.

## Bound 2: Memory Bound

$$\beta \geq \frac{Q}{T} = \frac{W/T}{W/Q} = \frac{P}{I}$$

Rearranging:

$$P \leq \beta \cdot I$$

## Combining the Bounds

$$P \leq \min(\pi, \beta \cdot I)$$

## Combining the Bounds

$$P \leq \min(\pi, \beta \cdot I)$$

Taking logarithms:

$$\log_2(P) \leq \min(\log_2(\pi), \log_2(\beta) + \log_2(I))$$

## Combining the Bounds

$$P \leq \min(\pi, \beta \cdot I)$$


Taking logarithms:

$$\log_2(P) \leq \min(\log_2(\pi), \log_2(\beta) + \log_2(I))$$


This is a **piecewise linear function** in log-log space:

- For low  $I$ :  $P = \beta \cdot I$  (diagonal line, slope 1)
- For high  $I$ :  $P = \pi$  (horizontal line)
- **Ridge point:**  $I^* = \pi/\beta$

# The Roofline Plot




## Reading the Roofline



**Goal:** Move points up toward the roof, or right to increase  $I$ .

## Optimization 1: ILP (Instruction-Level Parallelism)


**Scenario:** Compute-bound code with a single accumulator (latency-limited).



**ILP does not change operational intensity** – same work, same data.

## Optimization 2: SIMD (Vectorization)


**Scenario:** Code that processes one element at a time.



**SIMD does not change operational intensity – same work, same data.**

## Optimization 3: Blocking (Tiling)

**Scenario:** Matrix multiplication – reduce data movement via cache reuse.



**Blocking reduces data movement  $Q$ , increasing operational intensity  $I = W/Q$ .**

## Computing Operational Intensity

---

## Example 1: Vector Addition (DAXPY)

$$y = \alpha x + y$$

```
// daxpy.cpp
for (int i = 0; i < n; i++)
    y[i] = alpha * x[i] + y[i];

$ cd examples && make run_daxpy
```

# Example 1: Vector Addition (DAXPY)

$$y = \alpha x + y$$

```
// daxpy.cpp
for (int i = 0; i < n; i++)
    y[i] = alpha * x[i] + y[i];

$ cd examples && make run_daxpy
```

Work:  $W = 2n$  flops (1 multiply + 1 add per element)

Data movement (cold cache):

- Read  $x$ :  $8n$  bytes
- Read  $y$ :  $8n$  bytes
- Write  $y$ :  $8n$  bytes

$$Q = 24n \text{ bytes}$$

# Example 1: Vector Addition (DAXPY)

$$y = \alpha x + y$$

```
// daxpy.cpp
for (int i = 0; i < n; i++)
    y[i] = alpha * x[i] + y[i];

$ cd examples && make run_daxpy
```

**Work:**  $W = 2n$  flops (1 multiply + 1 add per element)

**Data movement (cold cache):**

- Read  $x$ :  $8n$  bytes
- Read  $y$ :  $8n$  bytes
- Write  $y$ :  $8n$  bytes

$Q = 24n$  bytes

**Operational Intensity:**  $I = \frac{2n}{24n} = \frac{1}{12} \approx 0.083$  flops/byte

## Example 2: Dot Product

$$s = x^T y = \sum_{i=0}^{n-1} x_i \cdot y_i$$

```
// dot.cpp
double s = 0;
for (int i = 0; i < n; i++)
    s += x[i] * y[i];

$ cd examples && make run_dot
```

## Example 2: Dot Product

$$s = x^T y = \sum_{i=0}^{n-1} x_i \cdot y_i$$

```
// dot.cpp
double s = 0;
for (int i = 0; i < n; i++)
    s += x[i] * y[i];

$ cd examples && make run_dot
```

**Work:**  $W = 2n$  flops

**Data movement:**

- Read  $x$ :  $8n$  bytes
- Read  $y$ :  $8n$  bytes

$Q = 16n$  bytes

## Example 2: Dot Product

$$s = x^T y = \sum_{i=0}^{n-1} x_i \cdot y_i$$

```
// dot.cpp
double s = 0;
for (int i = 0; i < n; i++)
    s += x[i] * y[i];

$ cd examples && make run_dot
```

**Work:**  $W = 2n$  flops

**Data movement:**

- Read  $x$ :  $8n$  bytes
- Read  $y$ :  $8n$  bytes

$Q = 16n$  bytes

**Operational Intensity:**  $I = \frac{2n}{16n} = \frac{1}{8} = 0.125$  flops/byte

## Example 3: Matrix-Vector Multiplication

$$y = Ax$$

```
// gemv.cpp
for (int i = 0; i < n; i++) {
    y[i] = 0;
    for (int j = 0; j < n; j++)
        y[i] += A[i][j] * x[j];
}
$ cd examples && make run_gemv
```

## Example 3: Matrix-Vector Multiplication

$$y = Ax$$

```
// gemv.cpp
for (int i = 0; i < n; i++) {
    y[i] = 0;
    for (int j = 0; j < n; j++)
        y[i] += A[i][j] * x[j];
}
$ cd examples && make run_gemv
```

**Work:**  $W = 2n^2$  flops

**Data movement:**

- Read  $A$ :  $8n^2$  bytes
- Read  $x$ :  $8n$  bytes (reused  $n$  times if fits in cache)
- Write  $y$ :  $8n$  bytes

$Q \approx 8n^2$  bytes (dominated by  $A$ )

## Example 3: Matrix-Vector Multiplication

$$y = Ax$$

```
// gemv.cpp
for (int i = 0; i < n; i++) {
    y[i] = 0;
    for (int j = 0; j < n; j++)
        y[i] += A[i][j] * x[j];
}
$ cd examples && make run_gemv
```

**Work:**  $W = 2n^2$  flops

**Data movement:**

- Read  $A$ :  $8n^2$  bytes
- Read  $x$ :  $8n$  bytes (reused  $n$  times if fits in cache)
- Write  $y$ :  $8n$  bytes

$Q \approx 8n^2$  bytes (dominated by  $A$ )

**Operational Intensity:**

$$I = \frac{2n^2}{8n^2} = \frac{1}{4} = 0.25 \text{ flops/byte}$$

## Example 4: Matrix-Matrix Multiplication (Naive)

$$C = AB$$

```
// gemm.cpp - compares naive vs blocked
for (int i = 0; i < n; i++)
    for (int j = 0; j < n; j++)
        for (int k = 0; k < n; k++)
            C[i][j] += A[i][k] * B[k][j];

$ cd examples && make run_gemm
```

## Example 4: Matrix-Matrix Multiplication (Naive)

$$C = AB$$

```
// gemm.cpp - compares naive vs blocked
for (int i = 0; i < n; i++)
    for (int j = 0; j < n; j++)
        for (int k = 0; k < n; k++)
            C[i][j] += A[i][k] * B[k][j];

$ cd examples && make run_gemm
```

**Work:**  $W = 2n^3$  flops

**Data movement (naive, cache « matrices):**

- Each  $C_{ij}$  requires reading row of  $A$  and column of  $B$
- Total:  $Q \approx 8n^3$  bytes (terrible!)

## Example 4: Matrix-Matrix Multiplication (Naive)

$$C = AB$$

```
// gemm.cpp - compares naive vs blocked
for (int i = 0; i < n; i++)
    for (int j = 0; j < n; j++)
        for (int k = 0; k < n; k++)
            C[i][j] += A[i][k] * B[k][j];

$ cd examples && make run_gemm
```

**Work:**  $W = 2n^3$  flops

**Data movement (naive, cache « matrices):**

- Each  $C_{ij}$  requires reading row of  $A$  and column of  $B$
- Total:  $Q \approx 8n^3$  bytes (terrible!)

**Operational Intensity (naive):**

$$I = \frac{2n^3}{8n^3} = \frac{1}{4} \text{ flops/byte}$$

Same as matrix-vector!

## Example 4b: Matrix-Matrix Multiplication (Blocked)

With blocking (block size  $b$ ,  $3b^2 \leq \text{cache}$ ):

**Data movement:** We divide  $A$ ,  $B$ ,  $C$  into  $b \times b$  blocks. For each block  $C_{ij}$ :

$$C_{ij} = \sum_{k=1}^{n/b} A_{ik} B_{kj}$$

Each block multiply loads three  $b \times b$  blocks ( $A_{ik}$ ,  $B_{kj}$ ,  $C_{ij}$ ), fitting in cache since  $3b^2 \leq \gamma$ .

## Example 4b: Matrix-Matrix Multiplication (Blocked)

With blocking (block size  $b$ ,  $3b^2 \leq \text{cache}$ ):

**Data movement:** We divide  $A$ ,  $B$ ,  $C$  into  $b \times b$  blocks. For each block  $C_{ij}$ :

$$C_{ij} = \sum_{k=1}^{n/b} A_{ik} B_{kj}$$

Each block multiply loads three  $b \times b$  blocks ( $A_{ik}$ ,  $B_{kj}$ ,  $C_{ij}$ ), fitting in cache since  $3b^2 \leq \gamma$ .

**Total block multiplies:**  $(n/b)^3 = n^3/b^3$

**Bytes per block multiply:**  $3 \times 8b^2$  (but  $C_{ij}$  is reused across  $k$ )

Net loads: each of the  $n^2/b^2$  blocks of  $A$  is read  $n/b$  times, same for  $B$ :

$$Q = 2 \times \frac{n^2}{b^2} \times \frac{n}{b} \times 8b^2 = \frac{2n^3}{b} \times 8$$

## Example 4b: Matrix-Matrix Multiplication (Blocked)

With blocking (block size  $b$ ,  $3b^2 \leq \text{cache}$ ):

**Data movement:** We divide  $A$ ,  $B$ ,  $C$  into  $b \times b$  blocks. For each block  $C_{ij}$ :

$$C_{ij} = \sum_{k=1}^{n/b} A_{ik} B_{kj}$$

Each block multiply loads three  $b \times b$  blocks ( $A_{ik}$ ,  $B_{kj}$ ,  $C_{ij}$ ), fitting in cache since  $3b^2 \leq \gamma$ .

**Total block multiplies:**  $(n/b)^3 = n^3/b^3$

**Bytes per block multiply:**  $3 \times 8b^2$  (but  $C_{ij}$  is reused across  $k$ )

Net loads: each of the  $n^2/b^2$  blocks of  $A$  is read  $n/b$  times, same for  $B$ :

$$Q = 2 \times \frac{n^2}{b^2} \times \frac{n}{b} \times 8b^2 = \frac{2n^3}{b} \times 8$$

**Operational Intensity (blocked):**

$$I = \frac{W}{Q} = \frac{2n^3}{16n^3/b} = \frac{b}{8} \text{ flops/byte}$$

## Example 4b: Matrix-Matrix Multiplication (Blocked)

With blocking (block size  $b$ ,  $3b^2 \leq \text{cache}$ ):

**Data movement:** We divide  $A$ ,  $B$ ,  $C$  into  $b \times b$  blocks. For each block  $C_{ij}$ :

$$C_{ij} = \sum_{k=1}^{n/b} A_{ik} B_{kj}$$

Each block multiply loads three  $b \times b$  blocks ( $A_{ik}$ ,  $B_{kj}$ ,  $C_{ij}$ ), fitting in cache since  $3b^2 \leq \gamma$ .

**Total block multiplies:**  $(n/b)^3 = n^3/b^3$

**Bytes per block multiply:**  $3 \times 8b^2$  (but  $C_{ij}$  is reused across  $k$ )

Net loads: each of the  $n^2/b^2$  blocks of  $A$  is read  $n/b$  times, same for  $B$ :

$$Q = 2 \times \frac{n^2}{b^2} \times \frac{n}{b} \times 8b^2 = \frac{2n^3}{b} \times 8$$

**Operational Intensity (blocked):**

$$I = \frac{W}{Q} = \frac{2n^3}{16n^3/b} = \frac{b}{8} \text{ flops/byte}$$

## Summary: Operational Intensity

Operation	$W$	$Q$	$I$
DAXPY: $y = \alpha x + y$	$2n$	$24n$	$1/12$
Dot product: $x^T y$	$2n$	$16n$	$1/8$
GEMV: $y = Ax$	$2n^2$	$8n^2$	$1/4$
GEMM naive: $C = AB$	$2n^3$	$8n^3$	$1/4$
GEMM blocked: $C = AB$	$2n^3$	$16n^3/b$	$b/8 \sim O(\sqrt{\gamma})$

## Summary: Operational Intensity

Operation	$W$	$Q$	$I$
DAXPY: $y = \alpha x + y$	$2n$	$24n$	$1/12$
Dot product: $x^T y$	$2n$	$16n$	$1/8$
GEMV: $y = Ax$	$2n^2$	$8n^2$	$1/4$
GEMM naive: $C = AB$	$2n^3$	$8n^3$	$1/4$
GEMM blocked: $C = AB$	$2n^3$	$16n^3/b$	$b/8 \sim O(\sqrt{\gamma})$

**Key insight:** Only GEMM (with blocking) can achieve high operational intensity!

## Roofline Examples

---

## Example Platform: Intel Skylake

### Single core specifications:

- Peak:  $\pi = 16 \text{ flops/cycle}$  (2 FMA units  $\times$  4-wide AVX  $\times$  2)
- At 3 GHz:  $\pi = 48 \text{ GFLOPS}$

## Example Platform: Intel Skylake

### Single core specifications:

- Peak:  $\pi = 16 \text{ flops/cycle}$  ( $2 \text{ FMA units} \times 4\text{-wide AVX} \times 2$ )
- At 3 GHz:  $\pi = 48 \text{ GFLOPS}$

### Memory bandwidth:

- Let's assume:  $\beta \approx 15 \text{ GB/s} = 5 \text{ bytes/cycle}$  at 3 GHz

## Example Platform: Intel Skylake

### Single core specifications:

- Peak:  $\pi = 16 \text{ flops/cycle}$  (2 FMA units  $\times$  4-wide AVX  $\times$  2)
- At 3 GHz:  $\pi = 48 \text{ GFLOPS}$


### Memory bandwidth:

- Let's assume:  $\beta \approx 15 \text{ GB/s} = 5 \text{ bytes/cycle}$  at 3 GHz

### Ridge point:

$$I^* = \frac{\pi}{\beta} = \frac{16}{5} = 3.2 \text{ flops/byte}$$

# Skylake Roofline



## Where Do Common Operations Fall?

---

Memory Bound	Compute Bound
BLAS Level 1: DAXPY, Dot product, Norms	BLAS Level 3: GEMM, Matrix factorizations
BLAS Level 2: GEMV, Matrix-vector solve	Dense linear algebra
Stencils, SpMV, Graph algorithms	Convolutions, GEMM

---

## Numerical Example: DAXPY

**Problem:**  $y = 2x + y$ , vectors of length  $n = 10^8$

**Platform:** 3 GHz Skylake,  $\pi = 48$  GFLOPS,  $\beta = 15$  GB/s

# Numerical Example: DAXPY

**Problem:**  $y = 2x + y$ , vectors of length  $n = 10^8$

**Platform:** 3 GHz Skylake,  $\pi = 48$  GFLOPS,  $\beta = 15$  GB/s

**Analysis:**

- $W = 2 \times 10^8$  flops
- $Q = 24 \times 10^8$  bytes = 2.4 GB
- $I = 1/12 \approx 0.083$  flops/byte

# Numerical Example: DAXPY

**Problem:**  $y = 2x + y$ , vectors of length  $n = 10^8$

**Platform:** 3 GHz Skylake,  $\pi = 48$  GFLOPS,  $\beta = 15$  GB/s

**Analysis:**

- $W = 2 \times 10^8$  flops
- $Q = 24 \times 10^8$  bytes = 2.4 GB
- $I = 1/12 \approx 0.083$  flops/byte

**Maximum achievable performance:**

$$P = \min(\pi, \beta \cdot I) = \min(48, 15 \times 0.083) = \min(48, 1.25) = 1.25 \text{ GFLOPS}$$

# Numerical Example: DAXPY

**Problem:**  $y = 2x + y$ , vectors of length  $n = 10^8$

**Platform:** 3 GHz Skylake,  $\pi = 48$  GFLOPS,  $\beta = 15$  GB/s

**Analysis:**

- $W = 2 \times 10^8$  flops
- $Q = 24 \times 10^8$  bytes = 2.4 GB
- $I = 1/12 \approx 0.083$  flops/byte

**Maximum achievable performance:**

$$P = \min(\pi, \beta \cdot I) = \min(48, 15 \times 0.083) = \min(48, 1.25) = 1.25 \text{ GFLOPS}$$

**Minimum runtime:**


$$T = \frac{W}{P} = \frac{2 \times 10^8}{1.25 \times 10^9} = 0.16 \text{ seconds}$$

## Adding More Roofs

---

## Multiple Compute Ceilings

Not all code can achieve peak  $\pi$ :




## Example: Instruction Mix Ceilings

**Your code's ceiling depends on:**

- Whether it vectorizes (SIMD)
- Whether compiler generates FMA
- Mix of adds vs multiplies


# Multiple Memory Ceilings

Different memory levels have different bandwidths:



**Implication:** Same code, different problem sizes = different rooflines!

## Effect of SIMD on Roofline



**SIMD increases  $\pi$ , shifting the ridge point right! \*\***

## Example: Matrix Multiplication - Batch Size Effect

Forward pass of fully-connected layer:  $Y = XW$  where  $X$  is batch  $\times$  input,  $W$  is input  $\times$  output

## Example: Matrix Multiplication - Batch Size Effect

Forward pass of fully-connected layer:  $Y = XW$  where  $X$  is batch  $\times$  input,  $W$  is input  $\times$  output

**Work:**  $W = 2 \times \text{batch} \times \text{input} \times \text{output}$

**Data:**

- $W$  matrix:  $8 \times \text{input} \times \text{output}$  bytes
- $X$  matrix:  $8 \times \text{batch} \times \text{input}$  bytes
- $Y$  matrix:  $8 \times \text{batch} \times \text{output}$  bytes

## Example: Matrix Multiplication - Batch Size Effect

Forward pass of fully-connected layer:  $Y = XW$  where  $X$  is batch  $\times$  input,  $W$  is input  $\times$  output

**Work:**  $W = 2 \times \text{batch} \times \text{input} \times \text{output}$

**Data:**

- $W$  matrix:  $8 \times \text{input} \times \text{output}$  bytes
- $X$  matrix:  $8 \times \text{batch} \times \text{input}$  bytes
- $Y$  matrix:  $8 \times \text{batch} \times \text{output}$  bytes

**Operational intensity:**

$$I \approx \frac{2 \times \text{batch} \times \text{in} \times \text{out}}{8(\text{in} \times \text{out} + \text{batch} \times \text{in} + \text{batch} \times \text{out})}$$

# Example: Matrix Multiplication - Batch Size Effect

Forward pass of fully-connected layer:  $Y = XW$  where  $X$  is batch  $\times$  input,  $W$  is input  $\times$  output

**Work:**  $W = 2 \times \text{batch} \times \text{input} \times \text{output}$

**Data:**

- $W$  matrix:  $8 \times \text{input} \times \text{output}$  bytes
- $X$  matrix:  $8 \times \text{batch} \times \text{input}$  bytes
- $Y$  matrix:  $8 \times \text{batch} \times \text{output}$  bytes


**Operational intensity:**

$$I \approx \frac{2 \times \text{batch} \times \text{in} \times \text{out}}{8(\text{in} \times \text{out} + \text{batch} \times \text{in} + \text{batch} \times \text{out})}$$

**For large batch:**  $I \rightarrow \frac{\text{batch}}{4}$  (weights reused across batch)

**For batch = 1:**  $I \approx 1/4$  (memory bound!)

## Batch Size and Roofline




This is why LLM decoding (batch=1) is memory bound.

Batching as an optimization works because a larger batch leads to a higher performance ceiling (ie doing more makes it more efficient, so we better do more)

## Using the Roofline for Optimization

---

# Optimization Strategy



## If Memory Bound...

**Goal:** Increase operational intensity  $I$  or get closer to bandwidth roof.

## If Memory Bound...

**Goal:** Increase operational intensity  $I$  or get closer to bandwidth roof.

### Techniques:

1. **Blocking/Tiling:** Reuse data in cache
2. **Loop fusion:** Combine loops to reuse loaded data
3. **Data layout:** Improve spatial locality
4. **Prefetching:** Hide memory latency
5. **Compression:** Reduce data volume

## If Compute Bound...

**Goal:** Increase performance  $P$  toward compute roof.

## If Compute Bound...

**Goal:** Increase performance  $P$  toward compute roof.


**Techniques:**

1. **Vectorization (SIMD):** Process multiple elements per instruction
2. **ILP:** Multiple accumulators, loop unrolling
3. **FMA:** Fused multiply-add instructions
4. **Algorithm improvement:** Reduce total work
5. **Parallelization:** Use multiple cores


## **Roofline Variations**

---

# Roofline for Different Data Types



# Hierarchical Roofline



## **Summary**

---

## The Roofline Model

$$P \leq \min(\pi, \beta \cdot I)$$

# The Roofline Model

$$P \leq \min(\pi, \beta \cdot I)$$

**Two bounds:**

1. **Compute bound:**  $P \leq \pi$  (limited by FLOPS)
2. **Memory bound:**  $P \leq \beta I$  (limited by bandwidth)

# The Roofline Model

$$P \leq \min(\pi, \beta \cdot I)$$

**Two bounds:**

1. **Compute bound:**  $P \leq \pi$  (limited by FLOPS)
2. **Memory bound:**  $P \leq \beta I$  (limited by bandwidth)

**Ridge point:**  $I^* = \pi/\beta$  separates the two regions

## The Roofline Model

$$P \leq \min(\pi, \beta \cdot I)$$

**Two bounds:**

1. **Compute bound:**  $P \leq \pi$  (limited by FLOPS)
2. **Memory bound:**  $P \leq \beta I$  (limited by bandwidth)

**Ridge point:**  $I^* = \pi/\beta$  separates the two regions

**Key insight:** Operational intensity  $I = W/Q$  determines which bound applies!

## Key Takeaways

1. **Calculate  $I$  before optimizing** – know your bottleneck!
2. **Memory bound ( $I < I^*$ )**: Improve locality, increase reuse
3. **Compute bound ( $I > I^*$ )**: SIMD, ILP, better algorithms
4. **Multiple roofs** exist for different:
  - Instruction mixes (FMA vs non-FMA)
  - Cache levels (L1 vs L2 vs DRAM)
  - Data types (FP64 vs FP32 vs FP16)
5. **Measure** with perf counters to validate analysis