

## **LLM Inference System Characteristics**

---

## **LLM Inference: Overview**

---

## What is LLM Inference?

Given a prompt, an LLM predicts the **next token**:

"The capital of France is"

- Tokenizer → [token IDs]
- Embedding Lookup → [vectors]
- Transformer Layer 1 → ... → Layer 32
- LM Head → [logits over vocabulary]
- Sample → "Paris"

## What is LLM Inference?

Given a prompt, an LLM predicts the **next token**:

"The capital of France is"

- Tokenizer → [token IDs]
- Embedding Lookup → [vectors]
- Transformer Layer 1 → ... → Layer 32
- LM Head → [logits over vocabulary]
- Sample → "Paris"

**Key insight:** LLMs generate text **one token at a time** (autoregressive).

# Autoregressive Generation



Each generated token becomes input for the next. We cannot generate token  $n + 1$  until we've generated token  $n$ .

## Running Example: Llama-3.1-8B


Parameter	Value	Description
Layers	32	Number of transformer layers
Hidden dimension ( $d$ )	4,096	Size of token representations
Query heads	32	Number of attention heads for queries
KV heads	8	Number of KV heads (GQA)
Head dimension	128	Dimension per attention head
FFN dimension	14,336	Intermediate size in feed-forward network
Vocabulary size	128,256	Number of possible tokens

Total: ~8 billion parameters

## Model Architecture


---

## High-Level Dataflow



- **Embedding:** Token ID  $\rightarrow$  4,096-dim vector (pure memory lookup)
- **32 Transformer Layers:** Attention + FFN
- **LM Head:** Project to vocabulary size, sample next token

# Transformer Layer Structure




Two blocks per layer:

1. **Attention block:** RMSNorm → Self-Attention → Residual add
2. **FFN block:** RMSNorm → SwiGLU FFN → Residual add

# Attention Mechanism

**Self-attention** allows each token to gather information from all previous tokens:

1. **Projections:**  $Q = xW_Q$ ,  $K = xW_K$ ,  $V = xW_V$
2. **Scores:**  $\text{softmax}(QK^\top / \sqrt{d_h})$
3. **Output:** Weighted sum of values, then  $W_O$  projection



## Multi-Head Attention Variants

**The KV cache problem:** During generation, we cache K and V to avoid recomputation.

Variant	Q Heads	KV Heads	KV Cache Size	Quality
MHA	32	32	1× (baseline)	Best
MQA	32	1	1/32×	Lower
<b>GQA</b>	32	8	1/4×	Good


**Llama-3.1-8B uses GQA:** Each of 8 KV heads is shared by 4 query heads.

⇒ 4× less KV cache memory, ~same quality

# Feed-Forward Network (SwiGLU)

$$\text{FFN}(x) = W_{\text{down}} \cdot (\text{SiLU}(W_{\text{gate}} \cdot x) \odot W_{\text{up}} \cdot x)$$

- $W_{\text{gate}}, W_{\text{up}}$ :  $4096 \rightarrow 14336$
- $W_{\text{down}}$ :  $14336 \rightarrow 4096$
- SiLU:  $z \cdot \sigma(z)$
- $\odot$ : element-wise multiply (gating)



## Local vs. Global Operations

**Key insight:** Attention is the *only* operation where tokens interact!

Operation	Scope	Description
Embedding	Per-token	Each token ID maps to its own vector
RMSNorm	Per-token	Normalizes each token independently
Q, K, V projections	Per-token	Matrix-vector multiply per token
FFN (SwiGLU)	Per-token	Same transformation for each token
<b>Attention</b>	<b>Global</b>	<b>Each token attends to all previous</b>

## Local vs. Global Operations

**Key insight:** Attention is the *only* operation where tokens interact!

Operation	Scope	Description
Embedding	Per-token	Each token ID maps to its own vector
RMSNorm	Per-token	Normalizes each token independently
Q, K, V projections	Per-token	Matrix-vector multiply per token
FFN (SwiGLU)	Per-token	Same transformation for each token
<b>Attention</b>	<b>Global</b>	<b>Each token attends to all previous</b>

Token-independent ops are embarrassingly parallel. Attention is what makes LLMs understand context.

## KV Cache, Prefill, and Decode

---

## The Redundancy Problem

When generating token  $t + 1$ , naive approach processes entire sequence:

$$[x_1, x_2, \dots, x_t, x_{t+1}]$$

## The Redundancy Problem

When generating token  $t + 1$ , naive approach processes entire sequence:

$$[x_1, x_2, \dots, x_t, x_{t+1}]$$

**But we already computed**  $K_1, K_2, \dots, K_t$  and  $V_1, V_2, \dots, V_t$  in previous steps!

## The Redundancy Problem


When generating token  $t + 1$ , naive approach processes entire sequence:

$$[x_1, x_2, \dots, x_t, x_{t+1}]$$

**But we already computed**  $K_1, K_2, \dots, K_t$  and  $V_1, V_2, \dots, V_t$  in previous steps!

**Solution:** Cache keys and values, reuse them.

# KV Cache



**KV cache** trades compute for memory:

- Store  $K$  and  $V$  for all previous tokens
- Each new token: compute only its  $K$ ,  $V$ , append to cache
- Attend using the full cache

## Two Phases of Inference

### Prefill (prompt processing)

- Process  $p$  prompt tokens in parallel
- Compute and cache all K, V
- Apply **causal mask** (token  $i$  only sees  $1..i$ )
- **Compute-bound** (matrix-matrix ops)

### Decode (token generation)

- Process 1 token at a time
- Compute new K, V; append to cache
- Attend to entire cache
- **Memory-bound** (matrix-vector ops)

Phase	Tokens/step	Bottleneck	Arithmetic Intensity
Prefill	$p$	Compute	High
Decode	1	Memory bandwidth	Low ( $\sim 1$ )

## KV Cache Memory Requirements

For Llama-3.1-8B with GQA:

$$\begin{aligned}\text{KV cache per token} &= 2 \times n_{\text{layers}} \times n_{\text{kv\_heads}} \times d_h \times \text{bytes} \\ &= 2 \times 32 \times 8 \times 128 \times 2 = 128 \text{ KB (FP16)}\end{aligned}$$

## KV Cache Memory Requirements

For Llama-3.1-8B with GQA:

$$\begin{aligned}\text{KV cache per token} &= 2 \times n_{\text{layers}} \times n_{\text{kv\_heads}} \times d_h \times \text{bytes} \\ &= 2 \times 32 \times 8 \times 128 \times 2 = 128 \text{ KB (FP16)}\end{aligned}$$

Context Length	KV Cache Size
4K tokens	0.5 GB
8K tokens	1 GB
32K tokens	4 GB
128K tokens	<b>16 GB</b>

At 128K context, KV cache equals model weight size!

## Causal Masking in Prefill

During prefill, we process all tokens in parallel but enforce causality:

$$\text{MaskedAttn}(Q, K, V) = \text{softmax} \left( \frac{QK^\top}{\sqrt{d_h}} + M \right) V$$

Causal Mask $M$					
$t_1$	0	$-\infty$	$-\infty$	$-\infty$	$-\infty$
$t_2$	0	0	$-\infty$	$-\infty$	$-\infty$
$t_3$	0	0	0	$-\infty$	$-\infty$
$t_4$	0	0	0	0	$-\infty$
$t_5$	0	0	0	0	0

$t_1 \quad t_2 \quad t_3 \quad t_4 \quad t_5$

After softmax:  $-\infty \rightarrow 0$

Token  $i$  only attends to  $1..i$

## Arithmetic Intensity Analysis

---

## Recap: Arithmetic Intensity

$$AI = \frac{\text{FLOPs}}{\text{Bytes transferred}}$$

**A100 GPU:** Peak = 312 TFLOPS / 2 TB/s = **156 FLOP/byte**

## Recap: Arithmetic Intensity

$$AI = \frac{\text{FLOPs}}{\text{Bytes transferred}}$$

**A100 GPU:** Peak = 312 TFLOPS / 2 TB/s = **156 FLOP/byte**

- AI > 156: **Compute-bound** (GPU busy computing)
- AI < 156: **Memory-bound** (GPU waiting for data)

## Linear Layers: Prefill vs Decode

For weight matrix  $W \in \mathbb{R}^{d_{in} \times d_{out}}$ :

Phase	Input Shape	FLOPs	Bytes (weights)	AI
Prefill	$(p, d_{in})$	$2 \cdot p \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	$\approx p$
Decode	$(1, d_{in})$	$2 \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	$\approx 1$

## Linear Layers: Prefill vs Decode


For weight matrix  $W \in \mathbb{R}^{d_{in} \times d_{out}}$ :

Phase	Input Shape	FLOPs	Bytes (weights)	AI
Prefill	$(p, d_{in})$	$2 \cdot p \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	$\approx p$
Decode	$(1, d_{in})$	$2 \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	$\approx 1$

**Prefill:** Amortize weight loading across  $p$  tokens  $\Rightarrow$  AI  $\approx p$

**Decode:** Each weight loaded once, used once  $\Rightarrow$  AI  $\approx 1$

## Decode is Memory-Bound



**Decode runs at <1% GPU utilization** for linear layers!

## Attention Arithmetic Intensity

**Prefill attention** (processing  $p$  tokens):

$$\text{AI}_{\text{prefill}} = \frac{p \cdot d}{d + d_{kv}} \approx 0.8p \quad (\text{for Llama-3.1-8B})$$

## Attention Arithmetic Intensity

**Prefill attention** (processing  $p$  tokens):

$$\text{AI}_{\text{prefill}} = \frac{p \cdot d}{d + d_{kv}} \approx 0.8p \quad (\text{for Llama-3.1-8B})$$

**Decode attention** (1 token,  $s$  cached tokens):

$$\text{AI}_{\text{decode}} = \frac{n_q}{n_{kv}} = \frac{32}{8} = 4 \quad (\text{GQA helps!})$$

## Attention Arithmetic Intensity

**Prefill attention** (processing  $p$  tokens):

$$\text{AI}_{\text{prefill}} = \frac{p \cdot d}{d + d_{kv}} \approx 0.8p \quad (\text{for Llama-3.1-8B})$$

**Decode attention** (1 token,  $s$  cached tokens):

$$\text{AI}_{\text{decode}} = \frac{n_q}{n_{kv}} = \frac{32}{8} = 4 \quad (\text{GQA helps!})$$

Compare to MHA: decode AI would be 1, not 4. **GQA improves decode efficiency.**

## Summary: Decode is the Bottleneck

Operation	Prefill AI	Decode AI
Linear layers	$\approx p$	$\approx 1$
Attention	$\approx 0.8p$	$\approx n_q/n_{kv}$
Normalization	$\approx 4$	$\approx 4$

With prompt length  $p = 1000$ : prefill is **compute-bound**.

At decode: **everything is memory-bound**.

## Summary: Decode is the Bottleneck

Operation	Prefill AI	Decode AI
Linear layers	$\approx p$	$\approx 1$
Attention	$\approx 0.8p$	$\approx n_q/n_{kv}$
Normalization	$\approx 4$	$\approx 4$

With prompt length  $p = 1000$ : prefill is **compute-bound**.

At decode: **everything is memory-bound**.

**The fundamental challenge:** Loading 16 GB of weights for *one* token's computation.

## Batching: The Key Optimization

---

## Why Batching Works

At batch size 1, each weight is loaded once, used once ( $AI \approx 1$ ).

**Batching:** Process  $B$  sequences simultaneously, reuse weights  $B$  times.

Batch Size	FLOPs	Bytes	AI
1	$2 \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	1
$B$	$B \cdot 2 \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	$B$

## Why Batching Works

At batch size 1, each weight is loaded once, used once ( $\text{AI} \approx 1$ ).

**Batching:** Process  $B$  sequences simultaneously, reuse weights  $B$  times.

Batch Size	FLOPs	Bytes	AI
1	$2 \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	1
$B$	$B \cdot 2 \cdot d_{in} \cdot d_{out}$	$2 \cdot d_{in} \cdot d_{out}$	$B$

**Arithmetic intensity scales linearly with batch size!**

To become compute-bound on A100: need  $B \approx 156$ .

## Attention Doesn't Batch Well

For attention, each sequence has its own KV cache:

Batch Size	FLOPs	Bytes (KV cache)	AI
1	$4 \cdot s \cdot d$	$4 \cdot s \cdot n_{kv} \cdot d_h$	$n_q/n_{kv}$
$B$	$B \cdot 4 \cdot s \cdot d$	$B \cdot 4 \cdot s \cdot n_{kv} \cdot d_h$	$n_q/n_{kv}$

## Attention Doesn't Batch Well


For attention, each sequence has its own KV cache:

Batch Size	FLOPs	Bytes (KV cache)	AI
1	$4 \cdot s \cdot d$	$4 \cdot s \cdot n_{kv} \cdot d_h$	$n_q/n_{kv}$
$B$	$B \cdot 4 \cdot s \cdot d$	$B \cdot 4 \cdot s \cdot n_{kv} \cdot d_h$	$n_q/n_{kv}$

**Batching does NOT improve attention's AI** — both compute and memory scale with  $B$ .

Exception: **prefix sharing** (shared system prompts) does help!

# The Batching Trade-off



- **Throughput** (tokens/sec total): improves with batch size
- **Latency** (time per request): barely increases at small batches!

At small  $B$ : GPU is memory-bound anyway—extra work is “free”.

## Memory Limits Batch Size

For Llama-3.1-8B at 4K context:

- Model weights: ~16 GB
- KV cache per sequence: ~0.5 GB
- On 80 GB GPU: room for ~128 concurrent sequences

## Memory Limits Batch Size

For Llama-3.1-8B at 4K context:

- Model weights: ~16 GB
- KV cache per sequence: ~0.5 GB
- On 80 GB GPU: room for ~128 concurrent sequences

**Practical batch sizes:** 32–64 sequences

**Key insight:** Batching is nearly universal in LLM serving because it improves throughput with minimal latency penalty.

# System Challenges from Batching

1. **Variable sequence lengths:** Requests have different prompt/output lengths
  - Solution: *Continuous batching* — dynamically add/remove requests
2. **Memory fragmentation:** KV caches grow unpredictably
  - Solution: *PagedAttention* — manage KV cache in fixed blocks
3. **Prefill-decode interference:** Different compute characteristics
  - Solution: *Chunked prefill* — interleave carefully

# System Challenges from Batching

1. **Variable sequence lengths:** Requests have different prompt/output lengths
  - Solution: *Continuous batching* — dynamically add/remove requests
2. **Memory fragmentation:** KV caches grow unpredictably
  - Solution: *PagedAttention* — manage KV cache in fixed blocks
3. **Prefill-decode interference:** Different compute characteristics
  - Solution: *Chunked prefill* — interleave carefully

These techniques form the core of vLLM, TensorRT-LLM, and SGLang.

## Architectural Variants

---

## Mixture of Experts (MoE)

Replace dense FFN with  $E$  expert FFNs; router selects top- $k$  per token:

$$\text{MoE}(x) = \sum_{i \in \text{TopK}(r(x))} g_i \cdot \text{FFN}_i(x)$$

Model	Total Params	Active Params	Experts	Top-K
Llama-3.1-8B	8B	8B	1 (dense)	—
Mixtral 8x7B	47B	13B	8	2
DeepSeek-V2	236B	21B	160	6

**MoE decouples total parameters from active parameters.**

# MoE Trade-offs

## Benefits:

- More parameters = better quality
- Same compute per token as smaller dense model
- 8× parameters, 2× compute (top-2)

## Challenges:

- All expert weights must fit in memory
- Load balancing: want even routing
- Communication in distributed settings

## Sliding Window Attention

Standard attention:  $O(n^2)$  — each token attends to all previous.

**Sliding window:** Limit attention to last  $w$  tokens:



- Mistral 7B:  $w = 4096$
- Bounds KV cache and compute regardless of sequence length
- Information propagates through layer stacking

## Multi-head Latent Attention (MLA)

**MLA** (DeepSeek-V2) compresses  $K$ ,  $V$  into low-dimensional latent:

$$c = W_{\text{compress}} \cdot x, \quad K = W_K \cdot c, \quad V = W_V \cdot c$$

Cache  $c$  instead of  $K$  and  $V$ .

Method	KV Cache per Token
MHA	$2 \cdot n_h \cdot d_h$
GQA	$2 \cdot n_{kv} \cdot d_h$ ( $4 \times$ less)
MLA	$d_c$ ( $\sim 5 \times$ less)

Trade-off: Extra compute to decompress during attention.

## **Summary**


---

## Key Takeaways

1. **LLMs generate tokens autoregressively** — sequential dependency limits parallelism
2. **Attention is the only global operation** — everything else is per-token
3. **Two phases:** Prefill (compute-bound, parallel) vs Decode (memory-bound, sequential)
4. **KV cache** trades compute for memory — critical for long contexts
5. **Decode is memory-bound** ( $AI \approx 1$ ) — loading weights dominates
6. **Batching** improves throughput by reusing weights across sequences

# The Fundamental Trade-off

LLM inference is dominated by memory bandwidth



**Data movement is the bottleneck** — just like CPUs, but at a different scale.